A Review of Modern Fashion Recommender Systems

YASHAR DELDJOO*, Polytechnic University of Bari, Italy
FATEMEH NAZARY, Polytechnic University of Bari, Italy
ARNAU RAMISA†, Amazon, USA
JULIAN MCAULEY, UC San Diego, USA
GIOVANNI PELLEGRINI, Polytechnic University of Bari, Italy
ALEJANDRO BELLOGIN, The Autonomous University of Madrid, Spain
TOMMASO DI NOIA, Polytechnic University of Bari, Italy

The textile and apparel industries have grown tremendously over the last years. Customers no longer have to visit many stores, stand in long queues, or try on garments in dressing rooms as millions of products are now available in online catalogs. However, given the plethora of options available, an effective recommendation system is necessary to properly sort, order, and communicate relevant product material or information to users. Effective fashion RS can have a noticeable impact on billions of customers' shopping experiences and increase sales and revenues on the provider-side.

The goal of this survey is to provide a review of recommender systems that operate in the specific vertical domain of garment and fashion products. We have identified the most pressing challenges in fashion RS research and created a taxonomy that categorizes the literature according to the objective they are trying to accomplish (e.g., item or outfit recommendation, size recommendation, explainability, among others) and type of side-information (users, items, context). We have also identified the most important evaluation goals and perspectives (outfit generation, outfit recommendation, pairing recommendation, and fill-in-the-blank outfit compatibility prediction) and the most commonly used datasets and evaluation metrics.

CCS Concepts: • **Information systems → Recommender systems**; **Multimedia and multimodal retrieval**; • **Human-centered computing → User models**.

Additional Key Words and Phrases: fashion recommender systems, visual, textile, size recommendation

---

*Corresponding author, Email: deldjooy@acm.org
†Work done outside Amazon employment

Authors' addresses: Yashar Deldjoo, yashar.deldjoo@poliba.it, Polytechnic University of Bari, Italy; Fatemeh Nazary, fatemeh.nazary@poliba.it, Polytechnic University of Bari, Italy; Arnau Ramisa, aramisay@amazon.com, Amazon, USA; Julian McAuley, jmcauley@eng.ucsd.edu, UC San Diego, USA; Giovanni Pellegrini, g.pellegrini4@studenti.poliba.it, Polytechnic University of Bari, Bari, Italy; Alejandro Bellogin, alejandro.bellogin@uam.es, The Autonomous University of Madrid, Madrid, Spain; Tommaso Di Noia, tommaso.dinoia@poliba.it, Polytechnic University of Bari, Italy.

111

1 INTRODUCTION

Recommender systems help users navigate large collections of products to find items relevant to their interests leveraging large amounts of product information and user signals like product views, followed or ignored items, purchases or web-page visits to determine how, when and what to recommend to their customers. Recommender systems have grown to be an essential part of all large Internet retailers, driving up to 35% of Amazon sales [103] or over 80% of the content watched on Netflix [31].

In this work we are interested in recommender systems that operate in one particular vertical market: garments and fashion products. This setting introduces a particular set of challenges and sub-problems, that are relevant for developing effective recommender systems.

Due to market dynamics and customer preferences, there is a large vocabulary of distinct fashion products, as well as high turnover. This leads to sparse purchase data, which challenges the usage of traditional recommender systems [83]. Furthermore, precise and detailed product information is often not available, making it difficult to establish similarity between products.

To deal with the aforementioned problems, and given the visual and aesthetic nature of fashion products, there is a growing body of computer vision research addressing tasks like localizing fashion items [50, 151], determining their category and attributes [24, 43, 145], or establishing the degree of similarity to other products [7, 46, 97, 99], to name only a few. Although works in the computer vision literature often don't consider personalization (or recommendation), their predictions and embeddings can be leveraged by recommender systems and combined with past user preferences, thus mitigating sparsity and cold start problems. Another relevant fashion problem that has attracted the attention of computer vision research is that of combining garments into complete outfits [17, 26, 48, 93]. Several works have studied how to learn the compatibility between fashion items using both professional photos of models wearing designer-created outfits, and social media pictures from 'influencers' and normal people [62, 127, 140]. In addition to allowing recommendations tailored to match the existing shopping basket or wardrobe of customers, these datasets help uncover other insights useful for recommender systems, such as the structure of fashion styles [61], social group preferences [81], or the evolution of trends across time and location [104, 106].

In addition to product-to-product relationships, fashion recommender systems also face particular product-to-user uncertainties, like *fit*, that can hurt the quality of recommendations if not taken into account. Fit prediction is a pain point for online fashion shopping according to customers, and the primary reason for product returns faced by online retailers [33]. The many sizing systems in use throughout the world, as well as their interpretation by different clothing manufacturers, make it very difficult to predict whether a particular product will properly fit a customer. Therefore, research on how to estimate a personalized garment fit has leveraged sources of information like co-purchase data [124], customer-reported measurements [112], or advanced imaging devices such as 3D scanners [11, 12, 35]. Needless to say, this information can be very valuable for recommendation.

Several online retailers, like Stitch-Fix and Amazon Prime Wardrobe, recreate the *try-out* experience of brick-and-mortar stores by shipping a highly tailored assortment of garments to a customer, who can proceed to try them on in their home, and return those that they do not like. Recommender systems can make this approach more sustainable by maximizing the number of items that customers keep [169], but the cold start problem and the commitment required from customers to stay in the membership plan make this approach hard to scale. Consequently, alternatives that would allow a customer to visualize the appearance and fit of clothes in an augmented reality or virtual environment are being researched [109].

The reverse approach, namely designing or automatically generating a computer model of a garment with the right appearance and fit for a customer and subsequently manufacturing it, is another route being explored by, for example, Amazon with the *Made for You Custom T-shirt* [1], with the aim of bringing the made-to-measure fashion to the masses.

1.1 Major Challenges

In this section we will describe the major challenges faced by recommender systems in the fashion domain.

Fashion item representation: Traditional recommender systems such as Collaborative Filtering or Content-Based Filtering have difficulties in the fashion domain due to the sparsity of purchase data, or the insufficient detail about the visual appearance of the product in category names [123]. Instead, more recent literature has leveraged models that capture a rich representation of fashion items through product images [51, 71], text descriptions or customer reviews [28, 164], or videos [153] which are often learned through surrogate tasks like classification or product retrieval. However, learning product representations from such input data requires large datasets to generalize well across different image (or text) styles, attribute variations, etc. Furthermore, constructing a representation that learns which product features customers take most into account when evaluating fashion products is still an open research problem.

Fashion item compatibility: Training a model that is able to predict if two fashion items 'go together,' or directly combine several products into an outfit, is a challenging task. Different item compatibility signals studied in recent literature include co-purchase data [107, 139], outfits composed by professional fashion designers [48], or combinations found by analyzing what people wear in social media pictures [72, 131]. From this compatibility information, associated image and text data is then used to learn to generalize to stylistically similar products. Some works explicitly model the latent style types [19]. An additional under-explored difficulty for compatibility prediction is the dependency on trends, seasonality, location or social group. Current approaches usually leverage image and text information.

Personalization and fit: The best fashion product to recommend depends on factors such as the location where the outfit will be used [8, 23, 80, 143], the season or occasion [98, 105, 162], or the cultural and social background of the customer [81, 131, 165]. A challenging task in fashion recommendation systems is how to discover and integrate these disparate factors [127, 147]. Current research often tackles these tasks by utilizing large-scale social media data.

As discussed earlier, a personalization dimension very particular to the fashion domain is that of fit. In addition to predicting what size of a product will be more comfortable to wear, body shape can influence stylistic choices [58, 60, 122].

Interepretability and Explanation: Most of the existing fashion recommender systems in the literature focus on improving predictive performance, treating the model as a black box. However, deploying accountable and interpretable systems able to explain their recommendations can foster user loyalty in the long term and improve the shopping experience. Current models generally offer explanations through highlighted image regions and attributes or keywords [47, 59, 89, 110, 149, 155].

Discovering Trends: Being able to forecast consumer preferences is valuable for fashion designers and retailers in order to optimize product-to-market fit, logistics and advertising. Many factors are confounded in what features are considered 'fashionable' or 'trendy', like seasonality [105], geographical influence [8], historical events [63] or style dynamics [9, 44, 100]. Again, social media is a useful resource leveraged by researchers [41, 67].

The challenges above, as well as many other issues have been discussed in the studied research works, and they are reviewed in the following sections. For instance, visual modeling of fashion items (cf. Section 2.2.2), modeling of fashion outfits (cf. Section 2.1.2), geo-temporal localization (cf. Section 2.2.3), attribute prediction for fashion and leveraging multi-modal data (cf. Section 2.2.2), explainability in fashion recommendation (cf. Section 2.1.4) and learning style (cf. Section 2.2.2).

1.2   Search strategy for relevant papers

We primarily relied on papers indexed in DBLP,[1] a prominent computer science bibliographic database, to identify publications that comprise the state of the art in fashion item and outfit recommendation. Our search approach was divided into two stages: finding relevant publication collections; and (ii) post-processing and filtering the final list. The total number of articles analyzed exceeds 50, with the majority of recognized research papers published between 2014 and 2021, an indication of the topic's originality and freshness. While we do not claim that this study is exhaustive in terms of the collected papers, we believe it gives a comprehensive overview of current achievements, trends, and what we deem to be the most significant challenges and tasks in fashion recommender systems. Additionally, it gives valuable information to both industry practitioners and academics.

1.3   Related Surveys

To put this survey in context, we identified and present related review and survey articles to explain in which ways our article differs from and extends earlier work. In a recent work, Deldjoo et al. [39] presented a survey of RS leveraging multimedia content, i.e., visual, audio, and/or textual features. The domains studied in this survey include various ones such as media streaming for audio and video recommendation, e-commerce for recommending different products including fashion items, news, and information recommendation, social media, and so forth. While fashion RS were also discussed, the authors only included a small portion of the topics and papers in this domain. Here, we discuss and present a comprehensive survey of significant tasks, challenges, and types of content used in the fashion RS field.

We have also identified surveys [29, 170] where the authors present a literature review of techniques at the intersection of fashion and computer vision (CV) and/or natural language processing (NLP). While we find these works relevant to this article, they remain largely different from the review presented here as those systems are not focused on RS but on other aspects of the fashion domain, such as text generation from images or pose estimation. Moreover, as another point of difference, we also provide recent techniques dealing with item visual and textual content representation exploited by RS approaches.

Perhaps the most relevant work to our current survey is a recent book chapter by Jaradat et al. [75] on fashion RS. This chapter focuses on discussing the state of the art of fashion recommendation systems; in particular, the authors affirm that deep learning represented a turning point with respect to the canonical approaches and therefore the authors examined four different tasks that use this new approach. Additionally they provided examples and possible problems and their evaluation. In particular, the authors focused their review on tasks related to social media and the size recommendation problem (see Section 2.1.3, where we introduce this task in detail). In our survey, in addition to analyzing the state of the art of the most commonly used algorithms in a wide range of tasks, we went in depth to understand which are the main features used by the more modern fashion recommender systems. In fact, an extensive discussion is held on how both the user and the items, with their characteristics, can be a source for the definition of models with accurate

---

[1]https://dblp.uni-trier.de/

Table 1. List of abbreviations used throughout this paper.

Abbreviation	Term	Abbreviation	Term
AI	Artificial Intelligence	G-RS	Graph-based RS
AML	Adversarial Machine Learning	GRU	Gated Recurrent Unit
AMR	Adversarial Multimedia Recommendation	IR	Information Retrieval
AUC	Area Under the ROC Curve	KG	Knowledge graph
BPR	Bayesian Personalized Ranking	K-NN	k-Nearest Neighbors
C&W	Carlini and Wagner	LFM	Latent Factor Model
CA	Context-Aware RS	LSTM	Long Short-Term Memory
CBF-RS	Content-Based Filtering RS	MF	Matrix Factorization
CF-RS	Collaborative Filtering RS	ML	Machine Learning
CNN	Convolutional Neural Network	NLP	Natural Language Processing
CS	Cold start	RNN	Recurrent Neural Network
CD-RS	Cross-Domain RS	ROC	Receiver Operating Characteristic
CV	Computer Vision	RS	Recommender Systems
DL	Deep Learning	SGD	Stochastic Gradient Descent
DNN	Deep Neural Network	SM	Social Media
ERM	Empirical Risk Minimization	SN	Social Network
FCN	Fully Convolutional Network	SVM	Support-Vector Machines
FGSM	Fast Gradient Sign Method	TF-IDF	Term Frequency–Inverse Document Frequency
FTF	Functional Tensor Factorization	URM	User Rating Matrix
GAN	Generative Adversarial Network	VRS	Visual recommender system

recommendations. Furthermore, since the fashion domain focuses on visual aspects, an in-depth study is also included here on computer vision approaches for increasing the understanding of item images.

2 CATEGORIZATION OF FASHION RECOMMENDER SYSTEMS

We categorize fashion recommender systems according to their task (cf. Section 2.1), and the input data they use to perform that task (cf. Section 2.2). We will discuss these two categorizations in more detail in the next sections.

2.1 Categorization based on task

By task we refer to the internal goal the fashion RS aims to achieve. This affects, in particular, the expected output of the algorithms. We identified five main tasks in the academic literature that fashion RS aim to achieve: *(i)* Fashion item recommendation, *(ii)* Fashion pair and outfit recommendation, *(iii)* Size recommendation, (iv) Explanation for Fashion Recommendation and *(v)* Other fashion prediction tasks. There are a few other tasks that so far have not gathered much attention, but are growing in popularity; they will be summarized later under the general subsection (v).

We discuss these three categories by presenting a broad definition of each task formally, while leaving a more detailed discussion of several prominent research works for Section 3

2.1.1 Fashion item recommendation: The fashion item recommendation task, similar to the classical recommendation problem, focuses on suggesting individual fashion items (clothing), that match users' preferences.

Definition 2.1 (Fashion item recommendation). Let $\mathcal{U}$ and $\mathcal{I}$ denote a set of users and fashion items in a system, respectively. Each user $u \in \mathcal{U}$ is related to $\mathcal{I}_u^+$, the set of items she has consumed.

Given a utility function $g : \mathcal{U} \times \mathcal{I} \rightarrow \mathbb{R}$ the **Item Recommendation Task** is defined as

$$\forall u \in \mathcal{U}, \ i_u^* = \underset{i \in \mathcal{I} \setminus \mathcal{I}_u^+}{\operatorname{argmax}} g(u, i) \tag{1}$$

where i_u^* denotes the best matching item not consumed by the user u before. The preference of user u on item i could be encoded as $s_{ui} \in \mathcal{S}$, a continuous-value score (e.g., 1-5 Likert scale), or implicit feedback in which we assume the user likes the item if she has interacted with (i.e., reviewed, purchased, clicked) the item (i.e., $s_{ui} = 1$). $\mathcal{I}_u^+$ represents the set of (u, i) pairs for which s_{ui} is known. The task of personalized Top-K fashion item recommendation problem is formally defined as identifying, for user u, a set of ranked list of items $X_u = \{i_1, i_2, ..., i_K\}$ that match user preference.

□

A simple yet effective pure CF model that serves as a foundation for many model-based VRS is BPR-MF [119]. Given a user u, and an item i, the core predictor in this model is given according to:

$$\hat{s}_{u,i} = p_u^T q_i \tag{2}$$

where $p_u, q_i \in \mathbb{R}^F$ are the embedding vectors for user u and item i, respectively, and F is the size of the embedding vector. This predictor is known as Matrix Factorization (MF), the parameters of the model can be learnt using Bayesian Personalized Ranking (BPR), a pairwise ranking optimization framework. We will show later in Section 3.1 how different authors have used this model as a starting point in the fashion domain, so that it is extended to consider other types of inputs and learning scenarios.

2.1.2 Fashion pair and outfit recommendation: Fashion outfits are sets of N items that are worn together, e.g., for an outdoor wedding, graduation party, baby shower, and so forth. The simplest form of a fashion outfit is when $N = 2$, i.e., two different items that look good when paired together, such as an orange shirt and a blue pair of jeans. However, in general outfits in online fashion store can be composed of items in different categories (e.g. show, bottom, top, hat, bag) that share some stylistic relationship.

Definition 2.2 (Fashion outfit composition). Let $O = \{i_1, i_2, ..., i_N\} \in O$ denote a fashion outfit composed of a set of compatible fashion items, in which $i_n \in \mathcal{I}$ is the n-th fashion item in the outfit O, and O is the set of all possible outfits, N is the length of the outfit, whose value is not fixed and can change with each outfit. Fashion Outfit Composition is formulated as follows: "given a scoring function $s(O)$ that indicates how well the outfit $O \in O$ is composed, find an outfit O that maximizes this utility":

$$O^* = \underset{O_j \in O}{\operatorname{argmax}} s(O_j) \tag{3}$$

where s is the outfit utility function, a scoring function that takes into account different types of relationships between fashion items to generate an outfit compatibility score. It is worth noting that, given this general task definition, the *evaluation* of fashion outfit composition is typically performed via fill-in-the-blank (FITB) or outfit compatibility score prediction [113] described in Section 4.

□

Moreover, it is possible to encode several objectives relevant to the fashion domain in the definition of the outfit composition scoring function by incorporating domain knowledge. For instance, McAuley et al. [107] defined compatibility according to *complementary* relationships (e.g., how much a white shirt complements blue pants), and *similarity* (how much one item in the outfit

is visually similar to another item). Hsiao and Grauman [62] model the outfit scoring function $s(O_j)$ as the superposition of two objectives:

$$O_j^* = \underset{O_j \in O}{\mathrm{argmax}}\, c(O_j) + v(O_j) \tag{4}$$

where $c(.)$ and $v(.)$ denote the *compatibility* score (how much pairs complement each other) and *versatility* score (defined as coverage of all styles), respectively. Parameters of the outfit scoring functions may be personalized to each style/user, and learned from user interaction data. As an example, in [93] the authors propose a system where compatibility is computed according to the image and category of every item in the outfit.

Recommending a pair of fashion items or an outfit evidences several challenges which have been addressed as follows:

- **Personalizing to a target customer.** Compatible fashion outfit recommendation can be studied in a personalized fashion with respect to that target user u to whom recommendations are computed:

$$O_u^* = \underset{O_j \in O}{\mathrm{argmax}}\, s(u, O_j). \tag{5}$$

A common problem here is the availability of sufficient interaction data to learn personalized outfit models. Different approaches have been proposed to address this issue in the literature, notably by using fashion generation [26], graph-learning approaches [87], binary codes [102], and finally self-attention mechanism to model the higher-order interactions between fashion items [101].

- **Modeling outfits as a sequence.** to take advantage of the representation of order-aware models such as LSTMs, some works [48] model an outfit as an ordered sequence of items, where each item belongs to step t. Then sequence-aware methods model the visual compatibility relationships of outfits, as a next item recommendation task:

$$\forall O_t = \{x_1, x_2, ..., x_t\}, \ x_{t+1} = \underset{x \in \mathcal{Y}}{\mathrm{argmax}}\, s(O_t, x) \tag{6}$$

where $\mathcal{Y}$ contains the set of allowed items. For instance, a user may want to find the best "shirt" that matches with his/her current "shoes" and "pants". In this example, all the "shirts" in the database could be a candidate item set, thus filling the cycle: "shoes → pants → shirts" and this cycle can go on to add more items such as "shoes → pants → shirts → hat → bag" (bottom-up outfit representation), depending on the size of outfit considered.

This problem might be generalized by allowing an input query x_q, either as a textual query (e.g., what outfit goes well with this mini skirt?) or a visual query (photo of a skirt). Moreover, when x_q is a visual query, the task is recognized as a **visual retrieval** task in the IR community or **contextual recommendation** in the RS community.

- **Content-based versus collaborative filtering.** while fashion item recommendation tasks have been dominated by CF-based approaches, which may or may not include side visual or textual information, the problem of outfit recommendation has been predominantly approached via content-based approaches, e.g., via visual metric learning problems. This can be explained due to scarcity of outfit interaction data. Nevertheless, there have been few works that approach the outfit problem in CBF-driven approach consider the work by Hu et al. [65], in which the authors use multi-modal features of items and tensor factorization to model the interactions between users and fashion items.

Here, we review some notable research works that study fashion outfit composition tasks. Two crucial issues arise here [158]:

- How can we learn domain knowledge about fashion compatibility relationships between fashion items in the presence of subtleties and subjectivity that can exist across customers?
- How can we incorporate the learned domain knowledge into the design space of fashion recommendation models?

A number of works have tried to model visual/style compatibility as a *metric learning* problem, for instance [54, 107, 113] for clothes and [6, 16] for furniture. Yin et al. [158] propose a fashion recommendation system that, in the first phase, learns and incorporates visual compatibility relationships of items at the pixel level. In the second phase, a domain adaptation strategy is used to allow the use of the knowledge extracted from the source domain that can be consistent with a different domain, that is, how to make sure that the knowledge learned in the database used in their experiments can also be compatible with others.

Polanía and Gupte [113] proposed a neural model composed of two sub-networks, first a Siamese sub-network which extracts visual features from a pair of images and a metric learning approach that maps the pair of features and auxiliary cues (color information) to a fashion compatibility score. Li et al. [88] proposed learning a representation of an outfit by exploiting the information from the image, title, and category and a multi-modal fusion model. Prediction is made by computing a compatibility score for the outfit representation. The scoring component is learned in an end-to-end fashion. Han et al. [48] model the outfit generation process as a sequential process and employ a bidirectional LSTM to predict the next item from the current one. However, these works build on two key assumptions: *(i)* a fixed number of items in an outfit or *(ii)* a fixed order of items in terms of their categories, which may not exist in reality. To address this issue, other approaches have been proposed that do not build on these assumptions. For instance, Lin et al. [90] address this issue by treating outfit recommendation as a multiple instance learning (MIL) problem, in which it is assumed that the labels for a group of items are known but not for all the individual items. The proposed system, named OutfitNet, is a two-stage framework for fashion outfit selection, where in the first step, a Fashion Item Relevancy network (FIR) learns the compatibility of fashion items and learns garment item relevance embeddings. In the second step, an Outfit Preference network (OP) uses visual input to learn the consumers' preferences for output. OutfitNet accepts various fashion products in a fashion outfit as input, learns the compatibility of the fashion items, the users' preferences for each item, and the users' attention on different pieces in the outfit.

Moreover, Chen et al. [26] presented an industrial scale fashion recommender system which they deployed on the Dida platform in Alibaba to provide personalized fashionable outfit recommendations to users. The main insight of the proposed approach, named personalized outfit generation (POG) is based on using a multi-modal neural model and a transformer architecture. The former creates a *multi-modal embedding* from several inputs of each item, dense CNN embedding plus dense text CNN of the item title, and dense CF embedding. The key advantage of the transformer is that it can better capture the dependencies among items in the user behavior sequence.

2.1.3 Size Recommendation. The capacity to take measures with your mobile phone and use algorithms to look at other more subjective factors like fit, the softness of the fabric, smell of the material, and so forth will become more and more important in the fashion recommendation domain Jaradat et al. [73].

A key customer concern that results in significant product returns is the *size-fit* problem. The choice of size is fundamental, as if incorrect the user will have to return the item creating economic damage to the company. Size recommendation systems have the primary purpose of helping users choose the items that best fit their size, without physically trying them on. According to [75], research on determining the right size and fit for customers faces the following challenges:

Lack of consistency between brands. There is a large number of approved sizing systems around the globe for various clothes, such as dresses, tops, skirts, pants and brands. Moreover, there are different size systems such as numeric (38-39-40), standard (S,M,L), fractions (41 1/3, 42.5), convention sizes (36-38, 40-42), country conventions (EU, FR, IT, UK), where inconsistencies and different ways of converting a local size system to another (as brands do not always comply with the same conversion logic) make the task challenging.

Subjectivity. The exact size is a very subjective feature; users who have purchased items with the same style and shape may make future purchases with different sizes; how an item fits on your body depends on or can be influenced by several factors, making an objective recommendation difficult. Moreover, customers may be driven by emotional aspects; even a piece of accurate size advice can come with a high emotional cost when the advised size differs from the customer's expectation.

Data sparsity. Users are able to buy only a small part of the items of an e-commerce website and on the other hand articles have a limited stock, which can in turn hinder the task of recommender systems working with user-item fit feedback.

Noise. The underlying interaction data where fashion recommenders are trained can be very noisy since users may make purchases not only for themselves but also share the information about clothing with their friends and family, hindering the task of size recommendation.

Given a fashion garment and the shopping history of a user, size recommendation methods predict whether a given size will be too large, small, or correct. Approaches for size recommendation can be classified according to which type of input data they rely on according to:

- *Physical body-related features.* The easiest way to make effective sizing recommendations is to use data from certain parts of the body [58, 60] such as bust, waist, and hip. Sometimes it may be useful to involve the user directly, for example by providing questionnaires to obtain other information, such as age, sex, height, and weight, or in more recent systems to provide their own photos and through vision algorithms to extract 3D avatars that reflect the physical characteristics of the users [82].
- *User-item fit feedback.* To provide personalized size recommendations, the interaction between the user and the item is essential; in this sense, using purchase data [2] or user returns and features of items [45, 58, 60] are used for this purpose, while allowing the user not to have to provide data relating to their physical appearance.

A few research works on size recommendation are reviewed here. Abdulla et al. [2] propose a size recommender system that treats the recommendation problem as a classification problem. It learns a joint size-fit shared latent space to project users and fashion items into, using the skip-gram-based Word2Vec model. A gradient boosting classifier is then used to predict the fit likelihood. The authors validate the usefulness of their system with both offline and online evaluation.

Guigourès et al. [45] propose a hierarchical Bayesian method for personalized size recommendation. The proposed system, which is dependent on customer and item pairings, provides the joint conditional probability of sizes ordered by consumers as well as their outcomes (i.e., article is kept, returned because it is too big, or returned because it is too small) as seen in training data. The approach makes personalized size suggestions by using the conditional probability of a size given a client and a fashion item. For parameter optimization and testing, the approach employs approximate probabilistic inference.

VIBE [60] (short for VIsual Body-aware Embedding) is a framework that captures clothing's association with various body shapes. Given an image of a user, the proposed embedding identifies clothing that will suit her body shape, where the embedding is learned from images of fashion models of various shapes and sizes wearing the garments. The proposed method allows having a
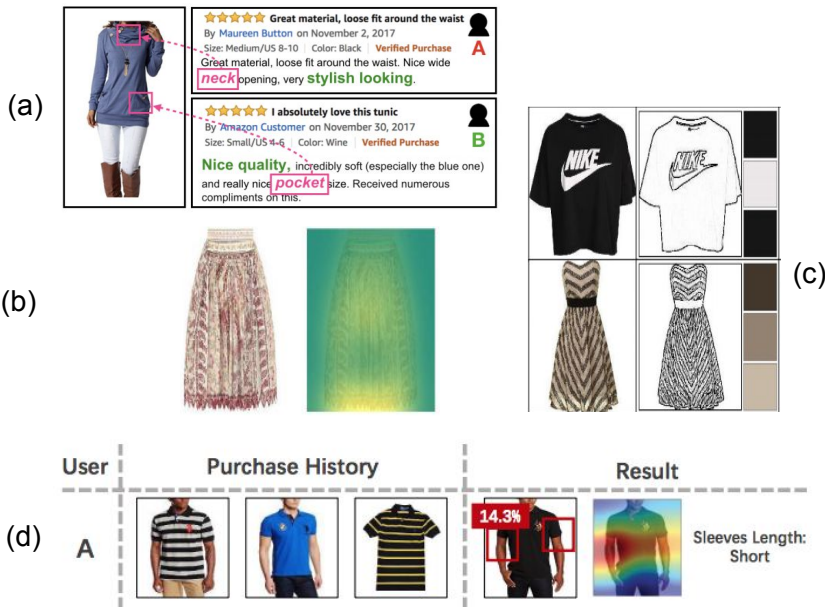
Fig. 1. (a) Personalized visual explanation by Chen et al. [28] related to region of the fashion image based on user reviews; (b) Explainable outfit recommendation by Lin et al. [92] using a neural attentive model; (c) outfit recommendation by Tangseng and Okatani [133], and (d) fashion item recommendation proposed by Hou et al. [59] comparing recommendation results with the target user's purchase history.

3D model of the human body that can be used for recommendations according to different body shapes.

In fashion, it is essential to know your body shape and chooses clothes to make your body look good. Hidayati et al. [58] suggested a clustering-based body shape (hourglass, rectangle, round, and inverted triangle) assignment approach, in which the body measurements of celebrities are compared to one another. Starting with photographs and information about celebrities, the system generates networks of images (by the visual appearance of clothing styles and body shape information) that propagate and automatically choose informative semantic traits; these are then used to propose clothing for each body shape. One issue with their approach is that the body shapes of celebrities are may not correspond to those of typical users

2.1.4 Learning explanations for fashion recommendation. An explanation is a piece of information that is shown to users that describes why a particular item is recommended. In recent years, there is a growing number of studies on *explainable recommendation* in different domains, notably in the fashion domain. The following dimensions are interesting to be studied:

- **Explanation goal.** Explanations have a significant impact on how individuals react to recommendations. A RS's aims in creating explanations for a particular recommendation, however, might range from enhancing the transparency to facilitating a speedier decision to convincing the user [13, 136].
- **Information source for explanation.** Recommendation explanations can be derived from many information sources and be presented in a variety of display styles. Examples of such styles in the fashion domain include textual explanations [28, 30], or visual explanations to

specific regions in the image [28], feature-level explanations based on fashion items' content features [36, 133], and by comparison with the user purchase history [59].

- **Explanation target:** Explainable recommendation research can target the explainability *recommendation methods* or the *recommendation results*. The former is concerned with devising interpretable models for increased transparency, while the latter treats the recommendation model as a black box, and develops a separate post-hoc/model-agnostic approach to explain recommendation results.

Table 2. Explanation goals and their definitions [13, 136]

Goal	Definition
Effectiveness	Help users make good decisions.
Efficiency	Help users make decisions more quicly.
Persuasiveness	Convince users to try or buy.
Satisfaction	Increase the ease of use or enjoyment.
Scrutability	Allow users to tell the system it is wrong.
Transparency	Explain how the system works.
Trust	Increase users' confidence in the system.

In Figure 1, different types of visual explanation used for fashion recommendation are illustrated. Chen et al. [28] present a method that provides visual explanations to the user through certain regions of the items, where an LSTM combines visual information and user reviews. With the same goal Tangseng and Okatani [133] propose a system that provides outfit recommendations and provides scores for each item or for each item feature in order to understand how much these influence the outfit composition. To explain whether an item/feature is compatible with the outfit and its influence on the score, three attributes have been extracted from the images: shape, texture, and colors. K-means clustering is applied for colors, and a CNN is used for shape and texture representation.

2.1.5 Other fashion-related prediction tasks. Besides fashion items and outfit recommendations, a few similar tasks are gaining attention in the fashion and retail industry. One of these areas includes *image-based fashion generation*, where methods have been proposed to generate outfits so that the user could explore them and judge their compatibility [79, 91, 156]. One of the key techniques used in these works is Generative Adversarial Networks (GANs), which will be described in detail in Section 3.2. Other relevant prediction tasks concern the fashion AI such as design [79], wardrobe creation [62], fashion trend forecasting [9], and societal/marketing biases such as socio-demographic inequality structures through recommendation systems in fashion [20], or user shopping behavior online and in brick-and-mortar stores [148], or capturing users' seasonal and temporal preference changes [150].

2.2 Categorization based on input

By input, we refer to the data fashion RS use to train their models. They could be based on a *U-I matrix* together with *side information beyond the U-I matrix* (related to users and item), and contextual information. This information determines the type of RS, according to CF, hybrid (CF+CBF) systems, and context-aware (CA) systems. In Table 3, we provide a classification of data/computational features used for the design of fashion RS by authors over the years.

Table 3. Classification of different research works incorporating various side information kinds in their proposed approach.

Authors	Year	Side-info User		Side information of items — Image representation								Side information of items — Text Representation								Other information			
		Social Network	Body features	Color	Texture	CNN	R-CNN	Siamese-CNN	RNN	Gramph-Embedding	Others	Pure	TF-IDF	CNN	RNN	Word2Vec	Transformer	Grapg-Embedding	Others	Social Influence	Fit Feedback	Pair Feedback	Others
Chen et al. [27]	2014					✓		✓						✓			✓			✓			
Polanía and Gupte [113]	2014							✓														✓	
Hou et al. [59]	2014					✓																	
Kumar and Gupta [84]	2014						✓															✓	
Cardoso et al. [21]	2014								✓			✓											
Tuinhof et al. [137]	2015					✓																	
Lin et al. [92]	2015					✓			✓													✓	
Heinz et al. [56]	2015				✓	✓			✓									✓					
Kang et al. [79]	2015					✓												✓				✓	
Bracher et al. [19]	2015					✓												✓					
Bhardwaj et al. [18]	2016			✓	✓																		✓
Abdulla et al. [2]	2017		✓													✓						✓	
Lin et al. [91]	2017					✓												✓					
Yin et al. [158]	2017					✓		✓													✓	✓	
Guigourès et al. [45]	2017		✓																✓				
Jaradat [72]	2017			✓				✓													✓		
Hidayati et al. [58]	2018		✓	✓		✓													✓		✓		
Agarwal et al. [4]	2018											✓											✓
Hwangbo et al. [69]	2018											✓											✓
Sun et al. [131]	2018					✓															✓	✓	
Zhou et al. [167]	2018				✓	✓		✓				✓											
Hu et al. [65]	2018		✓													✓							
Cardoso et al. [21]	2018				✓	✓										✓							
Jagadeesh et al. [71]	2018		✓	✓		✓																	
Liu et al. [94]	2018					✓	✓																
Zhang et al. [163]	2018		✓																				✓
Qian et al. [115]	2018	✓	✓	✓		✓		✓															✓
Zhao et al. [164]	2018								✓					✓									✓
McAuley et al. [107]	2018			✓	✓																		
Pan et al. [111]	2019								✓														
Sun et al. [131]	2019				✓	✓										✓							
Tangseng and Okatani [133]	2020			✓	✓						✓												
Sridevi et al. [130]	2020					✓																	
Brand and Gross [20]	2020													✓									✓
Xing et al. [150]	2020					✓		✓															✓
Wölbitsch et al. [148]	2020										✓												✓
Li et al. [86]	2020					✓						✓											
Banerjee et al. [15]	2020		✓			✓						✓											
Dong et al. [42]	2020										✓												✓
Ravi et al. [118]	2020				✓		✓																
Hidayati et al. [57]	2021		✓	✓		✓															✓	✓	
Yu et al. [160]	2021		✓			✓						✓											✓
Zhan et al. [161]	2021		✓						✓									✓					
Wen et al. [146]	2021					✓												✓					
Qiu et al. [116]	2021																	✓					

2.2.1 *Fashion RS relying on user-item interaction data.* Recommendation models that are trained purely from U-I interaction data are known as collaborative filtering (CF). Neighborhood models are among traditional CF classes that predict the unknown U-I preferences based on the neighborhoods by finding user-user similarities or item-item similarities, known as user-based or item-based CF, respectively, according to the preferences (or interactions). Machine-learned recommendation models are powered by model-based CF, a parameterized model whose parameters are learned

in the context of an optimization framework, such as Bayesian Personalized Ranking (BPR) [120]. Matrix factorization (MF), or BPR-MF, is one of the most promising approaches to building CF. MF essentially encodes the complex relations between users and items into a lower-dimensional (latent) representation of users and items, whose dot product explains the unknown predictions [83]. Thanks to deep neural models' ability to handle nonlinear data through nonlinear activation functions, recently *deep neural CF models* have been adopted by the research community to address the linearity issue of MF-based approaches and thereby uncover a more complex relationship between users and items. Given the challenges and nuances in the fashion domain, in the following, we review a few research works that enhance existing CF methods to incorporate the characteristics of fashion products. However, it should be noted that, for the inherent characteristics of this domain, most of the works incorporate additional or side information to the problem and, hence, will be discussed in following subsections.

Hwangbo et al. [69] propose a fashion RS that combines offline recommendation with online recommendation data. For a given fashion product that the customer favored before, the system considers two sets of items for the recommendation that the user potentially may want to buy from, products that (i) substitute and (ii) complement the earlier products. The authors further consider a preference decay factor to account for the life span of a product. The proposed system is named K-RecSys and is tested online. The evaluation reveals that purchases are higher for users in the recommendation group and that substitute recommendation is more frequently used compared to complementary recommendation.

The fashion catalog in e-commerce platforms is typically vast. Agarwal et al. [4] address this in two-stage recommendation formulation. In the first stage, the system proposes a similar product recommendation approach that takes a user query as input and generates a set of similar candidate products in a non-personalized fashion. In a second stage, a model-based CF using BPR provides user-level personalization (top 50 styles). Such approaches are suitable in the multimedia domain as caching a large number of items can be computationally cumbersome. Hu et al. [66] build an item-based CF model that enhances similarity calculation.

2.2.2 Fashion RS relying on user-item interaction and side-information. The range of information sources that RSs adopt can stretch beyond the U-I matrix, typically provided by attributes or information related to items such as category of user such as user's gender, age, and demographics [125]. To surface valuable recommendations, it is essential to learn from available user interactions, understand, and uncover the underlying decision factors. We categorize this side information according to items and users.

Fashion RS relying on U-I data and side-information of items. Understanding clothing provides a good platform for making recommendations [34]. State-of-the-art fashion RS employ a variety of types of content features, such as visual and textual features, as side-information to U-I interaction data. In the following, we review the most prominent features and attributes used for item representation in fashion recommendation literature [34, 39]:

(1) **Color.** The most common means to identify how one looks is achieved via colors, materials, and silhouettes on the body. Color and color consistency are among the most recognizable fashion clothing features that consumers often consider to decide what they want to wear. Color can be represented in several forms. For instance, as shown in Figure 2, Al-Rawi and Beel [10] describe a technique to segment and extract the color of a person's picture through probabilistic color modeling for clothing items. This technique can either increase or replace the data entry procedure needed to add a fashion item to electronic commerce catalogs. Jagadeesh et al. [71] evaluate two image representations, HSV Histograms and color bag-of-word (BoW). Their results showed that color is a better descriptor than textures.
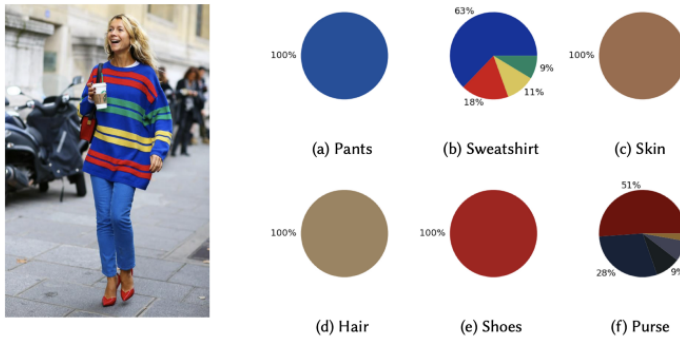
Fig. 2. Color extraction example from probabilistic color modelling of clothing items. Courtesy of [10].

In [115], a recommendation model is proposed that uses a color map through k-means based on the value of the segmented pixel of the clothing items with an aim toward extracting a dominant color. Starting from an attribute query, Chae et al. [22] convert this into a color vector space query; in particular, the Lab color space is used to simulate human visual perception through a palette of 269 colors.

(2) **Brand.** Product brands are a critical feature users consider when deciding among items. Wakita et al. [142] define a fashion-brand recommendation system, which through a deep (feed-forward) neural network has the primary purpose of predicting the user's favorite brand starting from user information, such as gender and height. Jaradat [72] propose a clothing recommendation system that uses brands as the main feature, where the model suggests to users clothes of similar brands that they have already purchased and that are on-trend.

(3) **Deep Visual Features.** Fashion recommendations can be improved using image-level features extracted through a deep network, such as a CNN [150, 167]. For instance, Zhou et al. [167] propose a CNN implementation to address the issue of two-piece apparel matching that is suitable with current fashion trends. They merged perception and reasoning models and constructed two parallel CNNs to enable the system to recognize garment features, one for upper-body clothing and another for lower-body apparel. A hierarchical topic model incorporates the resulting information into style topics with better semantic understanding to interpret the collocation patterns. The authors present a novel learning model based on Siamese Convolutional Neural Networks (SCNN) for learning a feature transformation from clothing photos to a latent feature space expressing fashion style consistency.

(4) **Texture.** The texture describes the body and surface of a garment. It has a direct impact on users, and is often used in recommendation systems. For example, Tangseng and Okatani [133] propose an outfit recommendation system based on attributes that are "human-interpretable," including texture and highlight how useful it is to use them. Zou et al. [171] use eight different Local binary patterns (LBPs) to extract texture features from the pixels of an input image, after which a local vector is calculated for each region and encoded to obtain the final vector. Qian et al. [115] provide a system that uses a CNN to locate textures in an image dataset to create matches and provide recommendations. On the other hand, Jagadeesh et al. [71] affirm that the color of the clothing is more important than the type of texture because the former is a stronger descriptor than the latter.

(5) **Style.** Style is associated with how people intend to express themselves through visual elements such as accessories, clothing, hair, and other aesthetic features. Measuring style computationally and offering personalized style-based fashion recommendations is important for retail companies and the central focus of many research works. A wide range of methods have been proposed to model style in fashion recommendation, ranging from modeling the body characteristics of people with the features of clothes presented in Hidayati et al. [57], modeling fashion substitution and complementarity with other products by learning a parametric transform of distances proposed by McAuley et al. [107] and multi-modal representations to build trend-sensitive models within the fashion field as proposed by Zhou et al. [167].

(6) **Knowledge graph (KG).** A KG is a heterogeneous graph where nodes serve as entities and edges serve as relationships between entities. The idea of making recommendations using side information from a KG has received considerable interest in the RecSys community. Such an approach can address the issues with recommender systems and provide explanations for recommended items. For instance, Li et al. [87] propose a model called a Hierarchical Fashion Graph Network, which uses a graph network to model the relationship between users and items for outfit recommendations through the propagation of information based on three levels, which are items, outfit level, and user level. Another possible use of KG is related to the solution of the cold start problem of many recommendation systems, for example, Yan et al. [152] use a model in which they build a user-item knowledge graph to discover the relationships that exist between them and solve the cold start problem. Zhan et al. [161] address the problem of personalised outfit recommendation by proposing an Attentive Attribute-Aware Fashion Knowledge Graph (A3-FKG) that is used to associate various outfits with both outfit- and item-level features.

(7) **Textual features.** Textual features in the fashion domain can have several forms e.g., reviews [28], article titles [164], article descriptions [91] and tags [19]. Zhao et al. [164] propose a *sentence model* that combines pairs of article titles to determine any compatibility of their styles. Textual features are integrated/modeled in the system for different purposes, for instance, for fashion generation [91]. In such a scenario, given an image of a top and a query vector description of the bottom, the fashion generator needs to generate a bottom image that matches the top image and the description as much as possible, process reviews [28] or attribute prediction [21]. A shown in Table 3, textual information has been used to represent item content in different forms, including TF-IDF [167], CNN [26, 164], Word2Vec [21, 65], transformer architecture [27], graph-embedding [45, 58] among others.

Some works have tried to evaluate the quality of the features mentioned above for fashion recommendation. For example, Jagadeesh et al. [71] compare textual and visual features or Deldjoo et al. [37] compare different CNN types for visual modeling of fashion items. In particular, the aim of [71] is to make efficient use of a large amount of fashion street images. To this end, two data-driven models are proposed, a deterministic fashion recommender (DFR) and a stochastic fashion recommender (SFR). The first is useful for identifying recommendations that are as objective as possible, the latter is used to model the bias of users, in fact it is known that the choice of a dress by a buyer is subjective. An important aspect highlighted is that color is more important than the type of texture, which underlines how a simple descriptor such as color can be important. Deldjoo et al. [37] evaluate the quality of several visual fashion recommender systems, including VBPR, DeepStyle, ACF, and VNPR (introduced in Section 3) empowered by pre-trained CNN types such as Alexnet, VGG-19, and ResNet 50. Evaluation is performed on accuracy and beyond-accuracy objective and qualitative assessment of the visual similarities between pairs of images. The results show that

ResNet50 overall provides relatively the best accuracy and beyond-accuracy recommendation performance over other pertained network types.

Fashion RS relying on interactions and side-information of users. These approaches in the most canonical form, use data from social networks and fashion magazines and website, in which user-user and body features (e.g., face, makeup, and size) information are extracted;

The main source of side information on users stretching beyond the U-I matrix is *social networks*, which has impacted a wide range of research disciplines in RS [125]. Mining social trust, influential figures, celebrities, and bloggers are among the influential factors impacting users' fashion choices. Different types of social network (SN)-based relationships have been used in the RS domain to enhance the quality of recommendation, including membership, friendship, following, and so forth.

Different types of online platforms have been used for the purpose of advertising, marking, social, and similar activities involving fashion and garments.

- **Social media.** With the advent of social media platforms like Instagram, the way fashion is advertised, and even the way it is designed is being re-imagined. Hundreds of thousands of people share images of their "outfit of the day," which prompts a flood of comments and queries from other users. Fashion brands can now use social media platforms such as Instagram to grow their businesses. Today, a single post by a fashion influencer receives much attention allowing many followers to know more about the brands and styles of the clothes that users mostly follow or to understand which brands and models users follow the most, or to catch the trends of the moment or seasonality [72, 74].

- **Fashion magazines and beauty websites.** It is especially common for celebrities' styles to be viewed as fashion inspirations as they pay fashion stylists to assist them with their wardrobe choices, which allows them to visually alter their actual body figure. In [58], the authors extract a list of essential female celebrities from prominent Fashion Magazines such as Vogue, Glamour, Marie Claire. In addition, a suite of body parameters is extracted from the body measurement website associated with celebrities. The dataset was completed using female celebrity photos taken from Google search Engines. Polyvore[2] was another good example of a fashion website, where users could build and post information about the fashion outfits reflecting their taste. These fashion outfits feature a wealth of multimodal data, including photographs and descriptions of fashion products, the number of times the outfit has been liked, and the outfit's hash tags. Researchers have applied this data to a variety of fashion-related tasks (cf. Section 4).

Sun et al. [131] propose a personalized clothing recommendation system, which takes into account both users' social circles and fashion style consistency of clothing products. Fashion style consistency refers to the fact that two clothing items, e.g., tops and skirts, can be visually different, but as long as they belong to the same style (e.g., sport, street, casual), the user may be likely buy them; hence fashion style consistency is an essential element for the design of clothing RS. To this end, the proposed approach considers three factors in the proposed clothing RS, namely interpersonal influence, personal interest, and interpersonal interest similarity; this was motivated by previous research works such as [114]. In particular, five types of matrices are built by mining the social data available and other sources of information: representing user-user social influence, representing the user-user similarity of interests, representing user-clothing similarity, representing clothing-clothing fashion style similarity, and representing user-clothing ratings. Afterward, a probabilistic matrix factorization (PMF) framework is used to integrate the above observed matrices and recommend suitable clothing products to users by casting the problem in an optimization

---

[2]www.polyvore.com

setting. Evaluation is carried out on real-world datasets collected from Moguije[3], a Chinese website for social fashion blending an SN with online shopping options.

2.2.3 Context-aware fashion RS. Context-aware (CA) models improve CF by including contextual information into the model to account for the current information need of the user [5, 39]. Context-awareness, i.e., understating the user's situational or contextual aspects, is essential in improving the user experience. Several taxonomies have been identified to classify context, e.g., based on computing environment, user environment, and physical environment or based on time, location, and social information (user's friends, social circles), see Deldjoo et al. [40] for more information. For the scope of this survey, we choose a pragmatic approach and organize CA Fashion RS according to the types of context frequently used in the fashion literature, namely *spatio-temporal context* (e.g., GPS coordinates, location, place of interest, time, season), *multimedia and physiological context* (e.g., free text, image or the user body/face features), and *affective context* (e.g., mood or emotion).

Spatio-temporal context. The *temporal* and *geographic* characteristics of fashion data influence human preferences. As for temporal aspects, first, fashion garments are sensitive to *seasonal* changes. Retail markets may face a decrease in preference over time after fashion products are released. Second, fashion *style trends* change as time progresses. For instance, trends can appear cyclically, with styles re-emerging after decades of absence, while others emerge and swiftly disappear. Understanding how style choices *vary with time*, and thereby what constitutes fashionable trends, is of crucial importance to the fashion retail business.

He and McAuley [52] propose an approach that involves learning the time evolution with which users decide to buy items through their implicit feedback, such as purchase history, clicks, and bookmarks. These features are used to define a predictor, helpful in approximating how much users interact with articles in a particular epoch. Dynamic FDNA proposed in [56] extends the static FDNA model with the inclusion of time-of-sale information. FDNA combines attributes and visual items using a neural architecture. The dynamic FDNA takes as input the timestamp of the purchase associated with a customer, and using existing information about the customer (purchase history so far), determines their current style. The dynamic model has the main advantage of providing recommendations for short-term customer intent. Adewumi et al. [3] propose a framework, which provides via a weather API recommendations for the exact day when the user needs them.

Multimedia and body/facial-related context. The term "multimedia context" refers to situations in which the system requires a user to provide a multimedia item (e.g., an image of a fashion item or an image plus textual query) in order to initiate the recommendation process (e.g., recommending a garment that complement an item the user is wearing).

- *Context = image.* Images are an important visual tool for users to communicate with a fashion recommender system. They could be used as input to the system to retrieve complementary or matching fashion products. For instance, given an image of a fashion item (e.g., "jeans"), one can identify matching fashion products (e.g., "tops") complement an item the user is considering [71, 167]. An application in the context of fashion is "upselling in e-commerce" when an online shopper with an item in her shopping basket is prompted to buy more things that match the one they already have in their shopping carts.
- *Context = image + text.* In addition to images, users may also include words (textual descriptions) to aid in the recommendation process. In such a scenario, the user shows a picture of an item (liked item) has and a description of what she is looking for (target item), and the model is responsible for considering both this information and recommending or generating matching

---

fashion items. For instance, Lin et al. [91] developed a model named FARM that allows the user to provide the image of the fashion item she likes (e.g., black top) and ask the system a query like "please suggest slim jeans". The model named FRAM decomposes the process into two parts, fashion recommender and fashion generation, for visual understanding and visual matching, respectively.

- *Context = body/facial-related attributes.* Context can also be attributed to user face or body shown to the system, where visual features extracted from the face or body can be used by the system to suggest matching cosmetic products [32, 163], clothing [57, 58] and hair-style recommendation [95, 154]. Hidayati et al. [57] propose a framework to provide size recommendations through both the characteristics of the body of the people and the style of the clothes. In [58] a system is presented in which the user should enter information about his or her body, such as shape, weight, height, bust, waist, hip this information is then correlated with photos of celebrities to define a style that can be close to that of the user.

In Table 4, we provide reference points to research works that explicitly use specific multimedia contexts (such as text and images) for fashion item information access (recommendation or search).

Table 4. Classification of fashion RS based on the input and output of the system. Compl. pair/outfit, refers to settings where **complementary** relationships between fashion items are considered, either for pairwise matches or outfit recommendation

Class	Input	Output	Research work
Classical	–	item	[45, 56, 59, 69, 72, 131, 163]
fashion RS	–	pair./outfit	[2, 27, 65, 69, 131, 133]
	image	item	[4, 18, 21, 79, 115, 137, 158, 166]
	image	pair/outfit	[18, 84, 92, 94, 107, 113, 131, 158, 164, 166]
Contextual	image+text	item	[167]
fashion RS	image+text	compl. pair/outfit	[91, 167]
	body/face	item	[32, 58]
	body/face	compl. pair/outfit	[163]
GAN-based	–	item	[79, 126, 128, 157]
fashion image gen.	–	compl. pair/outfit	[68, 84]

3 ALGORITHMS FOR FASHION RS

For the purpose of this survey, we focus on the two most prominent approaches for fashion RS, (i) visually-aware model-based CF (cf. Section 3.1), and (ii) generative fashion recommendation models (cf. Section 3.2).

3.1 Visually-aware model-based CF

(1) **VBPR (2016) [53].** The visual Bayesian Personalized Ranking (VBPR) method is built upon BPR and extends it by incorporating a latent content-based preference factor. The core predictor in VBPR is given by

$$\hat{x}_{u,i} = p_u^T q_i + \theta_u^T \theta_i \tag{7}$$

where $\theta_i \in \mathbb{R}^K$ is typically designed to represent the visual signal of an item, and $\theta_u \in \mathbb{R}^K$ represents the user preference on the visual dimension. He and McAuley [53] built the VBPR

predictor according to

$$\hat{x}_{u,i} = p_u^T q_i + \theta_u^T E \underbrace{\Phi_f(\mathbf{Img}_i)}_{f_i} \tag{8}$$

in which $f_i = \Phi_f(\mathbf{Img}_i) \in \mathbb{R}^{F \times 1}$ represents the latent feature of item i (typically extracted from a pretrained CNN) and Φ_f denotes the feature extractor. $E \in \mathbb{R}^{K \times D}$ linearly transforms the high-dimensional feature F into a lower-dimensional (e.g., $D = 20$) named the visual space. In [53], the authors show that VBPR improves the performance of BPR-MF and popularity-based recommenders for datasets chosen from Amazon fashion and Amazon phones.

(2) **DeepStyle (2017) [96].** VBPR, despite being effective, learns the visual information of items related to their category rather than their style (e.g., casual, aesthetic, formal). To build a dedicated style-CNN, one needs labeled data, and this could be a time-consuming process requiring expert fashion-related knowledge. The core idea explored by DeepStyle is to learn style from a U-I matrix. This was achieved by modeling `item = style + category`. Thus, to eliminate the category information, one can model `style = item - category`, transforming the VBPR predictor into

$$\hat{x}_{u,i} = p_u^T q_i + p_u^T (\mathbf{E} f_i - l_i) \tag{9}$$

where l_i is a latent factor that embodies the categorical information of item i.

(3) **ACF (2017) [25]** Traditional visual RS are not aware of the real meaning of users' interactions provided on multimedia content. When a user views several images on a social media platform, she might not have looked at all of these images with the same attention; similarly, she could have looked at different regions within the same image with varying importance levels. For instance, on a picture showing a dress over a sofa, she might like the picture because she likes dresses or chairs. Attentive collaborative filtering (CF) uses an attention network to incorporate two levels of attention, (i) item-level attention, and (ii) component-level attention (such as the regions in an image or frames of a video), in the core MR predictor according to

$$\hat{x}_{u,i} = \left(p_u + \sum_{k \in \mathcal{I}_u^+} a(u,k) v_k \right)^T q_i \tag{10}$$

in which $a(u,k)$ denotes user u's preference degree toward item k, and v_k is the attentive latent factor for item k, and q_i is the basic item vector in the MF model. Note that the proposed system is solved via a two-level neural architecture, which embodies two attention networks, $a(u,k)$ (presented above) and $\beta(u,k,l)$ denoting user u's preference degree in the l-th component of item k.

(4) **DVBPR (2017) [79]** The VBPR approach builds on CNNs (e.g., ResNet 50) to extract visual information. These networks are used in a pre-trained manner that (i) come from different domains (e.g., ResNet50 is trained on ImageNet), and second they have been trained to solve an image classification task and not a recommendation task). DVBPR is an end-to-end model, whose visual feature extractor is trained together with the whole preference prediction module (in a pair-wise manner).

$$\hat{x}_{u,i} = \theta_u^T \Phi_e(\mathbf{Img}_i) \tag{11}$$

in which $\mathbf{Img}_i$ is the image associated with product i and $\Phi_e(\mathbf{Img}_i)$ is the item content embedding where the CNN model is directly trained in a pair-wise manner.

Other relevant approaches here include TVBPR by He and McAuley [52], which takes into account seasonality and temporal changes into the VBPR model, and CO-BPR [158] for compatible outfit recommendation based on BPR.

Broadly speaking, the above techniques borrow ideas from more traditional forms of content-based recommendation, with the main challenge being how to handle the much higher dimensionality (and computational costs etc.) associated with visual data. Such data is handled either by projecting image representations into lower-dimensional spaces (as with VBPR), or through end-to-end (pixel level) representations.

3.2 Generative Fashion Recommendation Models

Most conventional RS are not suitable for application in the fashion domain due to unique characteristics hidden in this domain. Recently, GAN-based models have been used for fashion generation and fashion recommendation with promising performance. GANs gain their power exploiting their *generative power* allowing them to synthesize realistic-looking fashion items Deldjoo et al. [38]. This aspect can inspire the aesthetic appeal/curiosity of customers and designers and motivate them to explore the space of potential fashion styles.

[**CRAFT**] Huynh et al. [68] address the problem of recommending complementary fashion items based on visual features by using an adversarial process that resembles a GAN and uses a conditional feature transformer as $\mathcal{G}$ and a discriminator $\mathcal{D}$. One main distinction between this work and the prior literature is that the ⟨input, output⟩ pair for $\mathcal{G}$ are both features (here features are extracted using pre-trained CNNs [132]), instead of ⟨image, image⟩ or hybrid types such as ⟨image, features⟩ explored in numerous previous works [141, 168]. This would allow the network to learn the relationship between items directly on the feature space, spanned by the features extracted. The proposed system is named complementary recommendation using adversarial feature transform (CRAFT) since in the model, $\mathcal{G}$ acts like a feature transformer that—for a given query product image q—maps the source feature s_q into a complementary target feature $\hat{t}_q$ by playing a min-max game with $\mathcal{D}$ with the aim to classify fake/real features. For training, the system relies on learning the co-occurrence of item pairs in real images. In summary, the proposed method does not generate new images; instead it learns how to generate features of the complementary items conditioned on the query item.

[**DVBPR**] Deep visual Bayesian personalized ranking (DVBPR) [79] is one of the first works to exploit the *visual generative power of GANs* in the fashion recommendation domain. It aims to generate clothing images based on user preferences. Given a user and a fashion item category (e.g., tops, t-shirts, and shoes), the proposed system generates new images—i.e., clothing items—that are consistent with the user's preferences. The contributions of this work are two-fold: first, it builds an end-to-end learning framework based on the Siamese-CNN framework. Instead of using the features extracted in advance, it constructs an end-to-end system that turns out to improve the visual representation of images. Second, it uses a GAN-based framework to generate images that are consistent with the user's taste. Iteratively, $\mathcal{G}$ learns to generate a product image integrating a *user preference maximization objective*, while $\mathcal{D}$ tries to distinguish generated images from real ones. Generated images are quantitatively compared with real images using the preference score (mean objective value), inception score [121], and opposite SSIM [108]. This comparison shows an improvement in preference prediction in comparison with non-GAN based images. At the same time, the qualitative comparison demonstrates that the generated images are realistic and plausible, yet they are quite different from any images in the original dataset—they have standard shape and color profiles, but quite different styles.

[**MrCGAN**] Shih et al. [126] propose a compatibility learning framework that allows the user to visually explore candidate *compatible prototypes* (e.g., a white t-shirt and a pair of blue-jeans). The system uses metric-regularized conditional GAN (MrCGAN) to pursue the item generation task. It takes as input a projected prototype (i.e., the transformation of a query image in the latent "compatibility space"). It produces as the output a synthesized image of a compatible item (the

authors consider a compatibility notion based on the complementary of the query item across different catalog categories). Similar to the evaluation protocol in [68], the authors conduct online user surveys to evaluate whether their model could produce images that are perceived as compatible. The results show that MrCGAN can generate compatible and realistic images under compatibility learning settings compared to baselines.

[**Yang et al.** & c^+**GAN**] Yang et al. [157] address the same problem settings of MrCGAN [126] by proposing a fashion clothing framework composed of two parts: a clothing recommendation model based on BPR combined with visual features and a clothing complementary item generation based GAN. Notably, the generation component takes as input a piece of clothing recommended in the recommendation model and generates clothing images of other categories (i.e., tops, bottom, or shoes) to build a set of complementary items. The authors follow a similar qualitative and quantitative evaluation procedure as DVBPR [79] and further propose a *compatibility index* to measure the compatibility of the generated set of complementary items. A similar approach has also been proposed in c^+GAN [84], to generate a bottom fashion item paired with a given top.

3.3 Other Fashion Recommender System Algorithms

Depending on the task, fashion item and outfit recommendation, and size recommendation as we identified in Section 2.1, a variety of other models have been used in the studied research work. For example, RNNs [21, 92], graph-modeling [27, 87, 107], two-input (Siamese) CNNs [79, 113, 158, 164, 167], using attention mechanisms [92] and so forth. Detailed discussion on these approaches is left as a future direction.

3.4 Computer vision and fashion recommender systems

The fashion industry has a wide catalog of diverse items and a high rate of change or return as a result of market dynamics and customer preferences. This results in a lack of purchase data, which complicates the use of standard recommender systems. In addition, it is difficult to compute similarities across products when exact and extensive product information is not provided. Computer vision research is increasingly being used to address the issues outlined above. Given the visual and aesthetic nature of fashion products, a growing body of research addresses tasks like localizing fashion items, determining their category and attributes, or establishing the degree of similarity to other products.

A general classification of CV in Fashion RS is presented by Cheng et al. [29], according to (i) fashion detection (landmark detection, fashion parsing), (ii) fashion analysis (attribute prediction, style learning, popularity prediction), and (iii) fashion synthesis (style transfer, pose estimation). Fashion detection is a popular topic since most fashion CV techniques require performing fashion detection in the first place. For instance, Hsieh et al. [64] presents a visual try-on, a system that detects and generates clothing based on an image's human body parts. Fashion analysis is concerned with how people dress and the clothing they choose, as this can indicate their personalities and other social markers. Fashion synthesis is interested in visualizing what a person would look like with different cosmetics or garments by synthesizing realistic-looking images. CV techniques are also increasingly used in fashion RS. For instance, Qian et al. [115] use a CNN to segment and locate elements from the complicated backgrounds of street-style images. Hou et al. [59] use a semantic segmentation approach to derive region-specific attribute representations. CV techniques are also used in fashion generation for recommendation, e.g., Kumar and Gupta [84] use an enhanced conditional GAN to generate bottom items that can fit with the top items received in input by their framework.

4   EVALUATION AND DATASETS

Evaluation of fashion recommendations is a very complex task. It has to be adapted to the specific recommendation task (see Section 2.1) and consider the inputs of the algorithms (Section 2.2). Hence, in the next section we present several goals that aim to be achieved when deciding for a specific evaluation setting (Section 4.1), whereas Section 4.2 presents the most important perspectives analyzed in the literature; finally, Section 4.3 shows and discusses currently available datasets.

4.1   Evaluation goal

The goal of the evaluation is usually linked with the task the recommendation is intended to satisfy. However, we may find proposals where the algorithm is evaluated under different goals. For example, an outfit recommender can be evaluated only for that goal by ranking lists of outfits, but a more simple setting could be devised by using the recommender to score different outfits, in order to validate whether the algorithm is capable of discriminating those outfits that occur in the groundtruth.

In the following, we present the most important goals we have identified in the literature:

- **Goal 1 – Outfit generation**: The goal of this task is, given a fashion item x_q (e.g., skirt) representing the user's current interest, find the best item $x \in I$ (e.g., shirt) or the fashion outfit $F \in I$ (e.g., shirt, pants, hat) that goes/go well with the input query. The input x_q can be specified as a textual query (e.g., what outfit goes well to this mini skirt?) or a visual query (photo of a skirt).
  When x_q is a visual query, the task is recognized as a **visual retrieval** task in the community of information retrieval (IR) or **contextual recommendation** in the community of RS (x_q is a multimedia context representing a user's interest – see e.g. [78] for a definition of multimedia context (cf. Section 2.2.3).
- **Goal 2 – Outfit Recommendation**: This goal is related to the fashion and outfit recommendation task, where a set of objects is recommended to the user at once, by maximizing a utility function that measures the suitability of recommending a fashion outfit to a specific user.
- **Goal 3 – Pair recommendation**: This goal is a simplification of the outfit recommendation goal when $N = 2$. This is typically performed as *top-bottom* or *bottom-up* recommendation. In this task, given clothes related to the upper part of the body, the aim is to predict the possible lower part and vice versa. For instance, in a mobile setting, the user may take a photograph of a single piece of clothing of interest (e.g., tops) related to the upper part of the body and look for the bottom (e.g., trousers or skirts) from a large collection that best match the tops [71]. This typically requires a collection of pairs of top and bottom images for compatibility modeling. With advancements in computer vision, in some works, automated top-bottom region detection from photograph images has been proposed [70].
- **Goal 4 – Fill In the Blank** (FITB): This goal is used in a setting where we are given an incomplete outfit F^- (e.g., shirt, pants, accessories) with a missing item (e.g., shoes), and the method must find the best missing fashion item $x \in \{x_1, x_2, x_3, x_4\}$ from multiple choices where $F = \{F^-, x\}$ has fashion items, which are compatible visually [49]. This is a convenient scenario in real life, e.g., a user wants to choose a pair of shoes to suit his pants and jacket.
- **Goal 5 – Outfit compatibility prediction**: This goal is focused on, given a complete outfit $F \in \mathcal{F}$, predicting a compatibility score that best describes the composition of an outfit. All items in an outfit are compatible when all fashion items have similar style and go well together [49, 134]. This task is helpful since users may create their own outfits and wish to decide if they are compatible and trendy.

4.2 Evaluation perspective

Besides the goal that is considered when evaluating a recommender system, we may find authors that take different perspectives when considering what is a good recommendation. Although most of the literature in the fashion domain has been focused on building more accurate prediction and recommendation models, as in other recommendation settings, recent research has gone beyond this perspective and analyzed the explanations that should be presented to the user, the images that go with the recommendations, and even social aspects that may reinforce biases in the system. We describe these perspectives below.

- **Evaluate the recommendation**: Most papers perform offline evaluation, where classification or ranking metrics are very popular (precision, nDCG, AUC); in this sense, business metrics like CTR or purchase percentage are less common. See Table 6 for a summary. It should be noted that, depending on the recommendation task, other concepts such as compatibility might be measured to assess how well the recommendation is recognized by the user.
- **Evaluate an explanation**: Understanding recommendation explanations is generally a difficult task, since theoretically it can be evaluated only by real users. These approaches can be however costly and produce subjective judgements, however recent techniques aim to produce interpretable recommendations, instead of providing an explicit explanation in natural language. Examples of research works where this perspective is considered for the fashion domain are [149, 155].
- **Evaluate generated images**: When generating images, researchers typically measure three complementary metrics [79]: preference (how much each user would be interested in the recommended items), image quality using the inception score as a proxy, and diversity (where the visual similarity between every pair of returned images is computed). Sometimes, qualitative evaluation is also included. For example, Lin et al. [91] show several examples of generated items and discuss their validity. A similar qualitative evaluation is observed in [156].
- **Evaluate social perspectives**: Brand and Gross [20] highlights a very important aspect of today's society, gender equality, in detail the attention is placed on the price differences that emerge between men's and women's products. Other works like [117] argue whether users perceive in the same way recommendations generated by humans or by services exploiting Artificial Intelligence, in terms of trustworthiness and acceptance of the recommendations. In a similar line, the diversity of the recommendations (where the product exposure is higher) has been linked to a higher purchase rate and amount, although it may depend on the type of customer [85]. In any case, special care must be taken by automatic fashion recommender systems to not reinforce social biases and perpetuate long-observed objectifications [135, 138], hence, more research on these lines should be derived from the ML and RecSys communities.

From the summary presented in Table 6, we conclude that the most prominent perspective remains the first one (*evaluate the recommendation*). Nonetheless, some recent papers pay attention to other perspectives, such as image generation and social perspectives. This can be observed in the table by checking how many papers report image quality/diversity metrics, or metrics beyond accuracy, such as diversity or some kind of explanation measurements. The general trend, hence, is that offline evaluation with classical metrics (error, classification, and ranking) are prominent, but the community is working towards introducing others, more focused on alternative dimensions, such as business metrics or those related to social and explanation perspectives.

Table 5. The list of most commonly used datasets in the Fashion RS literature.

Dataset	Reference	Availability	Size	Area
Amazon Reviews	Link	Yes	140 million interactions	Generic
DeepFashion	[99]	No	8 hundred thousand images	Generic
Exact Street2Shop	Link	Yes	-	Street and shop photos
ExpFashion	[92]	No	200k images and 1 million comments	Item
FashionVC	[129]	No	20 thousand images	Outfits by fashion experts
Fashion-136K	[71]	No	-	Fashion models
Polyvore	[65]	No	-	Item photos
Pog	Link	Yes	4 millions images and descriptions	Item and Outfit
StyleReference	[57]	No	2 thousand pictures	Body features
Style4BodyShape	[58]	No	3150 celebrity records	Female celebrities

4.3 Available datasets

Table 5 summarizes a list of the most frequently used datasets, together with information on their availability, size, and the area to which the data pertain. They are discussed in more details in the following.

- **Amazon Reviews dataset.** Amazon is the world's largest marketplace for selling several categories of products, including fashion products. The collection consists of about 140 million records, including product reviews, ratings, and product information.
- **DeepFashion.** DeepFashion contains over 800K photos annotated with 50 categories, 1K attributes, and clothing landmarks (each image contains four to eight landmarks), as well as over 300K image pairs. They are classified according to several contexts, such as the store, street photo, and consumer.
- **Exact Street2Shop.** The images in this collection fall into two categories: *(a)* street images, which are photographs of individuals wearing various garments and taken in spontaneous, unplanned moments; *(b)* shop photos, photographs of clothing products from online clothing businesses, known as "shop photos".
- **ExpFashion.** This dataset was gathered from a collection of Polyvore images. They created new outfits given an item from current outfits and placed them in the dataset, starting with 1,000 seed outfits and expanding to a total of 100,000 outfits.
- **FashionVC.** Dataset consisting of the outfits of Polyvore's most fashion experts. Using a seed set of Polyvore popular outfits, the researchers identified 256 fashion experts, and then retrieved historical outfit from them.
- **Fashion-136K.** For this dataset, photographs of fashion models were gathered from the web. All photographs will have both top and bottom garments in the same image.
- **Polyvore.** A social commerce platform that allows users to exchange items and utilize them in image collages called Sets. As a dataset for many fashion recommendation systems (e.g., Hu et al. [65] have defined 150 sets of users), Sets are employed because they include only the products and a clean background, which substantially facilitates extraction of object properties.
- **POG.** The dataset contains data retrieved from the Taobao website; specifically, the clicks on the most popular items and outfits were extracted, and each one is paired with a record including the image's background, title, and category.
- **StyleReference.** There are 2,160 photos of apparel in this collection. The photographs were sourced from a number of major fashion publications style magazines (e.g., Elle and Vogue).
- **Style4BodyShape.** In this dataset, it is possible to find three different kinds of information: *(1)* a list of the most stylish 3,150 female celebrities, who are well-known for their refined

sense of fashion; *(2)* female celebrity body measurements, including dress, bust, waist, and hip measures, derived from a website that compiles this data, and *(3)* images of 270 fashionable celebrities crawled via Google search engine.

In summary, we notice that most of the reported datasets are not publicly available. This is a big hurdle to promote comparable and reproducible research in this domain. Moreover, we observe the area to which the data belongs to is heterogeneous, as not two datasets share the same area except the three datasets classified as *Generic*. Finally, even though most of these datasets are quite large (reaching millions of interactions), some of them are very small, including few thousands of images, which might be insufficient information for some methods to learn interesting patterns from them.

5 CONCLUSION AND FUTURE OUTLOOK

In this survey, we have analyzed and classified the RS that function in a specific vertical market: clothes and fashion goods. In particular, we have introduced a taxonomy of fashion recommender systems, which categorizes them according to the task (e.g., item, outfit, size recommendation, explainability among others), and type of side information (users, items, context). We have also identified the most important evaluation goals (outfit generation, outfit recommendation, pair recommendation, fill in the blank, and outfit compatibility prediction) and perspectives (evaluate the recommendation, the explanation, the generated images, or the social perspectives) exploited by the community, together with the most common datasets and evaluation metrics. This domain presents a unique collection of challenges and sub-problems pertinent to the development of successful recommender systems.

Datasets. We may recall how data collections began with simple "harvested" datasets and progressed to more "curated" datasets in subsequent years (e.g., images from fashion models rather than e.g., co-purchase data). It is interesting to consider which of these methods is preferable and where the field is headed in terms of datasets. Even though harvested datasets are easy to collect and correspond well to real prediction tasks (e.g., purchase estimate), collected datasets are noisy; they may also not reflect the real semantics of visual preferences, compatibility, or other aspects of a user's experience. In contrast, curated datasets may not match the distribution of real data; or "models" may not represent the preference dimensions of regular users, or the datasets may actually be contrary to one's goals. For instance, a marketing image may try to promote a pair of shoes by pairing it with a "boring" (non-distracting) outfit; thus curated data may not be any more "real" than harvested data (cf. Section 4.3).

Tasks. While traditional task in fashion RS research involved purchase or co-purchase prediction tasks, recent systems focus on combinatorial outputs (e.g. outfits and wardrobes) or even generative tasks (cf. Section 2.1). One needs to think whether this complexity is necessary. Do complicated non-pairwise functions actually exist in fashion semantics, or are pairwise functions adequate to represent the underlying decision factors? Is the increased complexity of combinatorial models worth the investment? This may require us to examine the long-term viability of simpler models in comparison to more complicated ones (cf. Section 2.1).

Interpretation. What exactly do visually aware models "learn"? Are they truly capturing fine-grained fashion semantics, or are they simply learning trivial factors (e.g., categories) from fashion data? (cf. Section 2.1.4) Visual recommender systems are mostly used to cope with noisy, sparse, cold-start datasets including visual data, but it is unclear that we have made much progress toward acquiring "real" fashion semantics in the manner of a designer. If our objective is just to predict purchases, this may be irrelevant; if our objective is to fulfill the function of fashion designers, how can we develop more representative datasets or tasks? (cf. Section 2.1.4)

Table 6. Common datasets and evaluation metrics used in the fashion recommender literature.

Authors	Year	type	Error	Clf.	Rank	Beyond	Expl.	Survey	business	image quality	image diversity	Clothing type	Datasets
Tangseng et al. [133]	2020	offline		✓			✓					generic	PO
Chen et al. [27]	2019	off+online		✓					✓			generic	POG
Polania et al. [113]	2019	offline			✓				✓			generic	PO
Hou et al. [59]	2019	offline		✓	✓							generic+shoe	Amazon
Kumar et al. [84]	2019	offline		✓	✓					✓	✓	top-bottom	Bing
Cardoso et al. [21]	2018	offline		✓								generic	ASOS
Tuinhof et al. [137]	2018	offline		✓								top-bottom	CrowdFlower
Lin et al. [92]	2020	offline		✓	✓		✓					top-bottom	FashionVC
Agarwal et al. [4]	2018	offline		✓	✓							Men T-Shirt	Myntra
He et al. [55]	2018	online									✓	generic	PO
Heinz et al. [56]	2017	offline		✓								generic	Zalando
Kang et al. [79]	2017	offline		✓				✓		✓	✓	shoe+top-bot.	Amazon
Bracher et al.. [19]	2016	online										generic	Zalando
Bhardwaj et al. [18]	2014	online						✓				top-bottom	Fashion-136K
Jagadeesh et al. [71]	2014	online						✓				top-bottom	Fashion-136K
Abdulla et al. [2]	2019	off+online										generic	NA
Lin et al. [91]	2019	offline		✓	✓			✓				top-bottom	ExpFashion
Yin et al. [158]	2019	offline		✓					✓			generic	Amazon
Guigoures et al. [45]	2018	offline		✓		✓						shoes	NA
Jaradat et al. [72]	2017	offline										-	Zalando+Instagram
Hidayati et al. [58]	2018	off+online			✓			✓				Woman Dress	Style4Body
Zhou et al. [167]	2018				✓							top-bottom	DeepFashion
Hu et al. [65]	2015	offline			✓							generic	PO
Cardoso et al. [21]	2018	offline		✓								generic	ASOS
Zhou et al. [166]	2019	online										generic	NA
Liu et al. [94]	2019	offline		✓								top-bottom	FarFetch
Hwangbo et al. [69]	2018	online							✓			generic	NA
Sun et al. [131]	2018	offline	✓	✓								top-bottom	MouJIE
Zhang et al. [163]	2017	offline		✓								-	NA
Wang et al. [144]	2015	online						✓				-	-
Qian et al. [115]	2017	offline		✓	✓							top-bottom	-
Zhao et al. [164]	2017	offline		✓								generic	Amazon+Tao
McAuley et al. [107]	2015	offline		✓								-	Amazon
Pan et al. [111]	2019	offline		✓				✓				generic	NA

Challenges IE: Interpretability/Explanation
Metrics Err.: Error-based accuracy (RMSE, MAE), Clf.: Classification metrics (Precision, Recall, F1, Accuracy, AUC)
RA: Rank-aware accuracy (MRR, MAP, NDCG), Beyond: Beyond accuracy metric (Novelty, Coverage, Diversity)
Expl.: Explanation related evaluation
Datasets PO: Polyvore, NA: Anonymous

Other aspects that we anticipate receiving future attention (and being the subject of future surveys) include the following:

Forecasting. The majority of fashion recommendations are focused on predicting the "present" based on previous interactions and trends, i.e., what the user will do *right now* based on their history. How might such models be utilized to make predictions regarding the fashion trends of the future? This aspect can be linked to popularity forecasting in the fashion domain since trendy items will likely be popular [14, 76].

Ethics. What are issues of fairness and bias in fashion recommendation, e.g. to what extent are underrepresented users (plus size, race, body type, etc.) poorly served by recommendations?

The ethics of "fast fashion" is also worth considering. Fast fashion is a term that refers to the manufacturing process, which means that the consumer may have to choose between purchasing something cheaper but manufactured under dubious conditions or paying more (or even having a very limited selection) for *sustainable* and *ethically* superior clothing.

Deployability. What are the practical considerations in terms of deploying models from academia? Perhaps complexity and being naturally black-box models (relating to the preceding discussion about interpretation) are important aspects for consideration given that the majority of the techniques described are based on neural networks.

Beyond these, fashion recommender systems are beginning to intersect with related domains including conversational models [159], models involving text, and augmented reality (virtual try on, etc.) [77].

REFERENCES

[1] [n.d.]. Amazon "Made for You" Custom T-shirt. https://www.amazon.com/Made-for-You-Custom-T-shirt/dp/B08N6J8G5M. Accessed: 2021-04-17.

[2] G. Mohammed Abdulla, Shreya Singh, and Sumit Borar. 2019. Shop your Right Size: A System for Recommending Sizes for Fashion products. In *Companion of The 2019 World Wide Web Conference, WWW 2019, San Francisco, CA, USA, May 13-17, 2019*. ACM, 327–334. https://doi.org/10.1145/3308560.3316599

[3] Adewole Adewumi, Adebola Taiwo, Sanjay Misra, Rytis Maskeliunas, Robertas Damasevicius, Ravin Ahuja, and Foluso Ayeni. 2019. *A Unified Framework for Outfit Design and Advice*. 31–41. https://doi.org/10.1007/978-981-13-9364-8_3

[4] Pankaj Agarwal, Sreekanth Vempati, and Sumit Borar. 2018. Personalizing Similar Product Recommendations in Fashion E-commerce. *CoRR* abs/1806.11371 (2018). arXiv:1806.11371 http://arxiv.org/abs/1806.11371

[5] Charu C Aggarwal. 2016. Ensemble-based and hybrid recommender systems. In *Recommender Systems*. Springer, 199–224.

[6] Divyansh Aggarwal, Elchin Valiyev, Fadime Sener, and Angela Yao. 2018. Learning Style Compatibility for Furniture. In *Pattern Recognition - 40th German Conference, GCPR 2018, Stuttgart, Germany, October 9-12, 2018, Proceedings (Lecture Notes in Computer Science)*, Thomas Brox, Andrés Bruhn, and Mario Fritz (Eds.), Vol. 11269. Springer, 552–566. https://doi.org/10.1007/978-3-030-12939-2_38

[7] Kenan E Ak, Ashraf A Kassim, Joo Hwee Lim, and Jo Yew Tham. 2018. Learning attribute representations with localization for flexible fashion search. In *Proceedings of the IEEE conference on computer vision and pattern recognition*. 7708–7717.

[8] Ziad Al-Halah and Kristen Grauman. 2020. From paris to berlin: Discovering fashion style influences around the world. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 10136–10145.

[9] Ziad Al-Halah, Rainer Stiefelhagen, and Kristen Grauman. 2017. Fashion forward: Forecasting visual style in fashion. In *Proceedings of the IEEE international conference on computer vision*. 388–397.

[10] Mohammed Al-Rawi and Joeran Beel. 2021. Probabilistic Color Modelling of Clothing Items. *Recommender Systems in Fashion and Retail* 734 (2021), 21.

[11] Phoebe R Apeagyei et al. 2010. Application of 3D body scanning technology to human measurement for clothing fit. *International Journal of Digital Content Technology and its Applications* 4, 7 (2010), 58–68.

[12] Susan P Ashdown, Suzanne Loker, Katherine Schoenfelder, and Lindsay Lyman-Clarke. 2004. Using 3D scans for fit analysis. *Journal of Textile and Apparel, Technology and Management* 4, 1 (2004), 1–12.

[13] Krisztian Balog and Filip Radlinski. 2020. Measuring Recommendation Explanation Quality: The Conflicting Goals of Explanations. In *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*. 329–338.

[14] Roja Bandari, Sitaram Asur, and Bernardo Huberman. 2021. The Pulse of News in Social Media: Forecasting Popularity. *Proceedings of the International AAAI Conference on Web and Social Media* 6, 1 (Aug. 2021), 26–33. https://ojs.aaai.org/index.php/ICWSM/article/view/14261

[15] Debopriyo Banerjee, Krothapalli Sreenivasa Rao, Shamik Sural, and Niloy Ganguly. 2020. BOXREC: Recommending a Box of Preferred Outfits in Online Shopping. *ACM Trans. Intell. Syst. Technol.* 11, 6 (2020), 69:1–69:28. https://doi.org/10.1145/3408890

[16] Sean Bell and Kavita Bala. 2015. Learning visual similarity for product design with convolutional neural networks. *ACM Trans. Graph.* 34, 4 (2015), 98:1–98:10. https://doi.org/10.1145/2766959

[17] Elaine M Bettaney, Stephen R Hardwick, Odysseas Zisimopoulos, and Benjamin Paul Chamberlain. 2019. Fashion outfit generation for E-commerce. *arXiv preprint arXiv:1904.00741* (2019).

[18] Anurag Bhardwaj, Vignesh Jagadeesh, Wei Di, Robinson Piramuthu, and Elizabeth F. Churchill. 2014. Enhancing Visual Fashion Recommendations with Users in the Loop. *CoRR* abs/1405.4013 (2014). arXiv:1405.4013 http://arxiv.org/abs/1405.4013

[19] Christian Bracher, Sebastian Heinz, and Roland Vollgraf. 2016. Fashion DNA: Merging Content and Sales Data for Recommendation and Article Mapping. *CoRR* abs/1609.02489 (2016).

[20] Alexander Brand and Tom Gross. 2020. Paying the Pink Tax on a Blue Dress - Exploring Gender-Based Price-Premiums in Fashion Recommendations. In *Human-Centered Software Engineering - 8th IFIP WG 13.2 International Working Conference, HCSE 2020, Eindhoven, The Netherlands, November 30 - December 2, 2020, Proceedings (Lecture Notes in Computer Science)*, Regina Bernhaupt, Carmelo Ardito, and Stefan Sauer (Eds.), Vol. 12481. Springer, 190–198. https://doi.org/10.1007/978-3-030-64266-2_12

[21] Ângelo Cardoso, Fabio Daolio, and Saúl Vargas. 2018. Product Characterisation towards Personalisation: Learning Attributes from Unstructured Data to Recommend Fashion Products. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining, KDD 2018, London, UK, August 19-23, 2018*, Yike Guo and Faisal Farooq (Eds.). ACM, 80–89. https://doi.org/10.1145/3219819.3219888

[22] Yeongnam Chae, Jiu Xu, Björn Stenger, and Soh Masuko. 2018. Color navigation by qualitative attributes for fashion recommendation. In *IEEE International Conference on Consumer Electronics, ICCE 2018, Las Vegas, NV, USA, January 12-14, 2018*. IEEE, 1–3. https://doi.org/10.1109/ICCE.2018.8326138

[23] Yu-Ting Chang, Wen-Haung Cheng, Bo Wu, and Kai-Lung Hua. 2017. Fashion world map: Understanding cities through streetwear fashion. In *Proceedings of the 25th ACM international conference on Multimedia*. 91–99.

[24] Huizhong Chen, Andrew Gallagher, and Bernd Girod. 2012. Describing clothing by semantic attributes. In *European conference on computer vision*. Springer, 609–623.

[25] Jingyuan Chen, Hanwang Zhang, Xiangnan He, Liqiang Nie, Wei Liu, and Tat-Seng Chua. 2017. Attentive Collaborative Filtering: Multimedia Recommendation with Item- and Component-Level Attention. In *SIGIR*. ACM, 335–344.

[26] Wen Chen, Pipei Huang, Jiaming Xu, Xin Guo, Cheng Guo, Fei Sun, Chao Li, Andreas Pfadler, Huan Zhao, and Binqiang Zhao. 2019. POG: personalized outfit generation for fashion recommendation at Alibaba iFashion. In *Proceedings of the 25th ACM SIGKDD international conference on knowledge discovery & data mining*. 2662–2670.

[27] Wen Chen, Pipei Huang, Jiaming Xu, Xin Guo, Cheng Guo, Fei Sun, Chao Li, Andreas Pfadler, Huan Zhao, and Binqiang Zhao. 2019. POG: Personalized Outfit Generation for Fashion Recommendation at Alibaba iFashion. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining, KDD 2019, Anchorage, AK, USA, August 4-8, 2019*. 2662–2670. https://doi.org/10.1145/3292500.3330652

[28] Xu Chen, Hanxiong Chen, Hongteng Xu, Yongfeng Zhang, Yixin Cao, Zheng Qin, and Hongyuan Zha. 2019. Personalized Fashion Recommendation with Visual Explanations based on Multimodal Attention Network: Towards Visually Explainable Recommendation. In *Proceedings of the 42nd International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2019, Paris, France, July 21-25, 2019*. ACM, 765–774. https://doi.org/10.1145/3331184.3331254

[29] Wen-Huang Cheng, Sijie Song, Chieh-Yun Chen, Shintami Chusnul Hidayati, and Jiaying Liu. 2021. Fashion meets computer vision: A survey. *ACM Computing Surveys (CSUR)* 54, 4 (2021), 1–41.

[30] Zhiyong Cheng, Xiaojun Chang, Lei Zhu, Rose C Kanjirathinkal, and Mohan Kankanhalli. 2019. MMALFM: Explainable recommendation by leveraging reviews and images. *ACM Transactions on Information Systems (TOIS)* 37, 2 (2019), 1–28.

[31] Sameer Chhabra. [n.d.]. Netflix says 80 percent of watched content is based on algorithmic recommendations. https://mobilesyrup.com/2017/08/22/80-percent-netflix-shows-discovered-recommendation/. Accessed: 2021-03-13.

[32] Kyung-Yong Chung. 2014. Effect of facial makeup style recommendation on visual sensibility. *Multim. Tools Appl.* 71, 2 (2014), 843–853. https://doi.org/10.1007/s11042-013-1355-6

[33] Jon Cilley. 2016. Apparel & footwear retail survey report. *Body Labs Inc, New York* (2016).

[34] Humberto Corona. [n.d.]. The State of Recommender Systems for Fashion in 2020. https://towardsdatascience.com/the-state-of-recommender-systems-for-fashion-in-2020-180b3ddb392f/. Accessed: 2021-13-19.

[35] Hein AM Daanen and Agnes Psikuta. 2018. 3D body scanning. In *Automation in garment manufacturing*. Elsevier, 237–252.

[36] Yashar Deldjoo, Tommaso Di Noia, Danielle Malitesta, and Felice Merra, Antonio. 2022. Leveraging Content-Style Item Representation for Visual Recommendation. In *Advances in Information Retrieval - 44th European Conference on IR Research, ECIR 2022*.

[37] Yashar Deldjoo, Tommaso Di Noia, Daniele Malitesta, and Felice Antonio Merra. 2021. A Study on the Relative Importance of Convolutional Neural Networks in Visually-Aware Recommender Systems. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 3961–3967.

[38] Yashar Deldjoo, Tommaso Di Noia, and Felice Antonio Merra. 2021. A survey on Adversarial Recommender Systems: from Attack/Defense strategies to Generative Adversarial Networks. *Comput. Surveys* (2021). https://doi.org/10.1145/

3439729

[39] Yashar Deldjoo, Markus Schedl, Paolo Cremonesi, and Gabriella Pasi. 2020. Recommender Systems Leveraging Multimedia Content. *ACM Comput. Surv.* 53, 5 (2020), 106:1–106:38. https://doi.org/10.1145/3407190

[40] Yashar Deldjoo, Markus Schedl, and Peter Knees. 2021. Content-driven Music Recommendation: Evolution, State of the Art, and Challenges. *arXiv preprint arXiv:2107.11803* (2021).

[41] Yujuan Ding, Yunshan Ma, Lizi Liao, Waikeung Wong, and Tat-Seng Chua. 2021. Leveraging Multiple Relations for Fashion Trend Forecasting Based on Social Media. *IEEE Transactions on Multimedia* (2021).

[42] Min Dong, Xianyi Zeng, Ludovic Koehl, and Junjie Zhang. 2020. An interactive knowledge-based recommender system for fashion product design in the big data environment. *Inf. Sci.* 540 (2020), 469–488. https://doi.org/10.1016/j.ins.2020.05.094

[43] Beatriz Quintino Ferreira, Joao P Costeira, Ricardo G Sousa, Liang-Yan Gui, and Joao P Gomes. 2019. Pose Guided Attention for Multi-Label Fashion Image Classification. In *2019 IEEE/CVF International Conference on Computer Vision Workshop (ICCVW)*. IEEE Computer Society, 3125–3128.

[44] Takao Furukawa, Chikako Miura, Kaoru Mori, Sou Uchida, and Makoto Hasegawa. 2019. Visualisation for analysing evolutionary dynamics of fashion trends. *International Journal of Fashion Design, Technology and Education* 12, 2 (2019), 247–259.

[45] Romain Guigourès, Yuen King Ho, Evgenii Koriagin, Abdul-Saboor Sheikh, Urs Bergmann, and Reza Shirvany. 2018. A hierarchical bayesian model for size recommendation in fashion. In *RecSys*. ACM, 392–396.

[46] M Hadi Kiapour, Xufeng Han, Svetlana Lazebnik, Alexander C Berg, and Tamara L Berg. 2015. Where to buy it: Matching street clothing photos in online shops. In *Proceedings of the IEEE international conference on computer vision*. 3343–3351.

[47] Xianjing Han, Xuemeng Song, Jianhua Yin, Yinglong Wang, and Liqiang Nie. 2019. Prototype-guided attribute-wise interpretable scheme for clothing matching. In *Proceedings of the 42nd International ACM SIGIR Conference on Research and Development in Information Retrieval*. 785–794.

[48] Xintong Han, Zuxuan Wu, Yu-Gang Jiang, and Larry S. Davis. 2017. Learning Fashion Compatibility with Bidirectional LSTMs. In *Proceedings of the 2017 ACM on Multimedia Conference, MM 2017, Mountain View, CA, USA, October 23-27, 2017*. ACM, 1078–1086. https://doi.org/10.1145/3123266.3123394

[49] Xintong Han, Zuxuan Wu, Yu-Gang Jiang, and Larry S Davis. 2017. Learning fashion compatibility with bidirectional lstms. In *Proceedings of the 25th ACM international conference on Multimedia*. ACM, 1078–1086.

[50] Kota Hara, Vignesh Jagadeesh, and Robinson Piramuthu. 2016. Fashion apparel detection: the role of deep convolutional neural network and pose-dependent priors. In *2016 IEEE Winter Conference on Applications of Computer Vision (WACV)*. IEEE, 1–9.

[51] Ruining He and Julian McAuley. 2016. VBPR: visual bayesian personalized ranking from implicit feedback. In *Thirtieth AAAI Conference on Artificial Intelligence*.

[52] Ruining He and Julian J. McAuley. 2016. Ups and Downs: Modeling the Visual Evolution of Fashion Trends with One-Class Collaborative Filtering. In *Proceedings of the 25th International Conference on World Wide Web, WWW 2016, Montreal, Canada, April 11 - 15, 2016*, Jacqueline Bourdeau, Jim Hendler, Roger Nkambou, Ian Horrocks, and Ben Y. Zhao (Eds.). ACM, 507–517. https://doi.org/10.1145/2872427.2883037

[53] Ruining He and Julian J. McAuley. 2016. VBPR: Visual Bayesian Personalized Ranking from Implicit Feedback. In *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence, February 12-17, 2016, Phoenix, Arizona, USA*, Dale Schuurmans and Michael P. Wellman (Eds.). AAAI Press, 144–150. http://www.aaai.org/ocs/index.php/AAAI/AAAI16/paper/view/11914

[54] Ruining He, Charles Packer, and Julian J. McAuley. 2016. Learning Compatibility Across Categories for Heterogeneous Item Recommendation. In *IEEE 16th International Conference on Data Mining, ICDM 2016, December 12-15, 2016, Barcelona, Spain*, Francesco Bonchi, Josep Domingo-Ferrer, Ricardo Baeza-Yates, Zhi-Hua Zhou, and Xindong Wu (Eds.). IEEE Computer Society, 937–942. https://doi.org/10.1109/ICDM.2016.0116

[55] Tong He and Yang Hu. 2018. FashionNet: Personalized Outfit Recommendation with Deep Neural Network. *CoRR* abs/1810.02443 (2018). arXiv:1810.02443 http://arxiv.org/abs/1810.02443

[56] Sebastian Heinz, Christian Bracher, and Roland Vollgraf. 2017. An LSTM-Based Dynamic Customer Model for Fashion Recommendation. In *RecTemp@RecSys (CEUR Workshop Proceedings)*, Vol. 1922. CEUR-WS.org, 45–49.

[57] Shintami Chusnul Hidayati, Ting Wei Goh, Ji-Sheng Gary Chan, Cheng-Chun Hsu, John See, Lai-Kuan Wong, Kai-Lung Hua, Yu Tsao, and Wen-Huang Cheng. 2021. Dress With Style: Learning Style From Joint Deep Embedding of Clothing Styles and Body Shapes. *IEEE Trans. Multim.* 23 (2021), 365–377. https://doi.org/10.1109/TMM.2020.2980195

[58] Shintami Chusnul Hidayati, Cheng-Chun Hsu, Yu-Ting Chang, Kai-Lung Hua, Jianlong Fu, and Wen-Huang Cheng. 2018. What Dress Fits Me Best?: Fashion Recommendation on the Clothing Style for Personal Body Shape. In *ACM Multimedia*. ACM, 438–446.

[59] Min Hou, Le Wu, Enhong Chen, Zhi Li, Vincent W. Zheng, and Qi Liu. 2019. Explainable Fashion Recommendation: A Semantic Attribute Region Guided Approach. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence, IJCAI 2019, Macao, China, August 10-16, 2019*. ijcai.org, 4681–4688. https://doi.org/10.24963/ijcai.2019/650

[60] Wei-Lin Hsiao and Kristen Grauman. 2020. ViBE: Dressing for Diverse Body Shapes. In *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2020, Seattle, WA, USA, June 13-19, 2020*. Computer Vision Foundation / IEEE, 11056–11066. https://doi.org/10.1109/CVPR42600.2020.01107

[61] Wei-Lin Hsiao and Kristen Grauman. 2017. Learning the latent" look": Unsupervised discovery of a style-coherent embedding from fashion images. In *Proceedings of the IEEE International Conference on Computer Vision*. 4203–4212.

[62] Wei-Lin Hsiao and Kristen Grauman. 2018. Creating capsule wardrobes from fashion images. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 7161–7170.

[63] Wei-Lin Hsiao and Kristen Grauman. 2021. From culture to clothing: Discovering the world events behind a century of fashion images. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*. 1066–1075.

[64] Chia-Wei Hsieh, Chieh-Yun Chen, Chien-Lung Chou, Hong-Han Shuai, Jiaying Liu, and Wen-Huang Cheng. 2019. FashionOn: Semantic-guided image-based virtual try-on with detailed human and clothing information. In *Proceedings of the 27th ACM International Conference on Multimedia*. 275–283.

[65] Yang Hu, Xi Yi, and Larry S. Davis. 2015. Collaborative Fashion Recommendation: A Functional Tensor Factorization Approach. In *ACM Multimedia*. ACM, 129–138.

[66] Zhi-Hua Hu, Xiang Li, Chen Wei, and Hong-Lei Zhou. 2019. Examining collaborative filtering algorithms for clothing recommendation in e-commerce. *Textile Research Journal* 89, 14 (2019), 2821–2835.

[67] Fu-Hsien Huang, Hsin-Min Lu, and Yao-Wen Hsu. 2021. From Street Photos to Fashion Trends: Leveraging User-Provided Noisy Labels for Fashion Understanding. *IEEE Access* 9 (2021), 49189–49205.

[68] Cong Phuoc Huynh, Arridhana Ciptadi, Ambrish Tyagi, and Amit Agrawal. 2018. CRAFT: Complementary Recommendation by Adversarial Feature Transform. In *ECCV Workshops (3) (Lecture Notes in Computer Science)*, Vol. 11131. Springer, 54–66.

[69] Hyunwoo Hwangbo, Yang Sok Kim, and Kyung Jin Cha. 2018. Recommendation system development for fashion retail e-commerce. *Electron. Commer. Res. Appl.* 28 (2018), 94–101. https://doi.org/10.1016/j.elerap.2018.01.012

[70] Tomoharu Iwata, Shinji Watanabe, and Hiroshi Sawada. 2011. Fashion coordinates recommender system using photographs from fashion magazines. In *Twenty-Second International Joint Conference on Artificial Intelligence*.

[71] Vignesh Jagadeesh, Robinson Piramuthu, Anurag Bhardwaj, Wei Di, and Neel Sundaresan. 2014. Large scale visual recommendations from street fashion images. In *The 20th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, KDD '14, New York, NY, USA - August 24 - 27, 2014*. ACM, 1925–1934. https://doi.org/10.1145/2623330.2623332

[72] Shatha Jaradat. 2017. Deep Cross-Domain Fashion Recommendation. In *Proceedings of the Eleventh ACM Conference on Recommender Systems, RecSys 2017, Como, Italy, August 27-31, 2017*. ACM, 407–410. https://doi.org/10.1145/3109859.3109861

[73] Shatha Jaradat, Nima Dokoohaki, Humberto Jesús Corona Pampín, and Reza Shirvany. 2020. Second Workshop on Recommender Systems in Fashion–fashionXrecsys2020. In *Fourteenth ACM Conference on Recommender Systems*. 632–634.

[74] Shatha Jaradat, Nima Dokoohaki, Kim Hammar, Ummul Wara, and Mihhail Matskin. 2018. Dynamic CNN models for fashion recommendation in Instagram. In *2018 IEEE Intl Conf on Parallel & Distributed Processing with Applications, Ubiquitous Computing & Communications, Big Data & Cloud Computing, Social Computing & Networking, Sustainable Computing & Communications (ISPA/IUCC/BDCloud/SocialCom/SustainCom)*. IEEE, 1144–1151.

[75] Shatha Jaradat, Guy Shani, Humberto Jesus, Corona Pampın, and Reza Shirvany. 2022. Fashion Recommender Systems. In *Recommender Systems Handbook*. Springer.

[76] Amin Javari and Mahdi Jalili. 2014. Accurate and Novel Recommendations: An Algorithm Based on Popularity Forecasting. *ACM Trans. Intell. Syst. Technol.* 5, 4, Article 56 (dec 2014), 20 pages. https://doi.org/10.1145/2668107

[77] Chamodi Jayamini, Asangika Sandamini, Thisara Pannala, Prabhash Kumarasinghe, Dushani Perera, and Kasun Karunanayaka. 2021. The use of augmented reality to deliver enhanced user experiences in fashion industry. *Lecture Notes in Computer Science* 12936 (2021).

[78] Marius Kaminskas and Francesco Ricci. 2012. Contextual music information retrieval and recommendation: State of the art and challenges. *Computer Science Review* 6, 2-3 (2012), 89–119. https://doi.org/10.1016/j.cosrev.2012.04.002

[79] Wang-Cheng Kang, Chen Fang, Zhaowen Wang, and Julian J. McAuley. 2017. Visually-Aware Fashion Recommendation and Design with Generative Image Models. In *2017 IEEE International Conference on Data Mining, ICDM 2017, New Orleans, LA, USA, November 18-21, 2017*. 207–216. https://doi.org/10.1109/ICDM.2017.30

[80] Hirokatsu Kataoka, Yutaka Satoh, Kaori Abe, Munetaka Minoguchi, and Akio Nakamura. 2019. Ten-million-order human database for world-wide fashion culture analysis. In *Proceedings of the IEEE/CVF Conference on Computer*

Vision and Pattern Recognition Workshops. 0–0.

[81] M Hadi Kiapour, Kota Yamaguchi, Alexander C Berg, and Tamara L Berg. 2014. Hipster wars: Discovering elements of fashion styles. In *European conference on computer vision*. Springer, 472–488.

[82] Daejin Kim, Wooyeol Ryu, Seungho Lee, Byungsu Lim, and Sangjun Lee. 2016. A Unity3D-based mobile fashion coordination system. *International Journal of Advanced Media and Communication* 6, 1 (2016), 86–92.

[83] Yehuda Koren, Robert M. Bell, and Chris Volinsky. 2009. Matrix Factorization Techniques for Recommender Systems. *IEEE Computer* 42, 8 (2009), 30–37. https://doi.org/10.1109/MC.2009.263

[84] Sudhir Kumar and Mithun Das Gupta. 2019. c$^+$GAN: Complementary Fashion Item Recommendation. *KDD '19, Workshop on AI for fashion, Anchorage, Alaska-USA* (2019).

[85] Hyokmin Kwon, Jaeho Han, and Kyungsik Han. 2020. ART (Attractive Recommendation Tailor): How the Diversity of Product Recommendations Affects Customer Purchase Preference in Fashion Industry?. In *Proceedings of the 29th ACM International Conference on Information & Knowledge Management* (Virtual Event, Ireland) *(CIKM '20)*. Association for Computing Machinery, New York, NY, USA, 2573–2580. https://doi.org/10.1145/3340531.3412687

[86] Eileen Li, Eric Kim, Andrew Zhai, Josh Beal, and Kunlong Gu. 2020. Bootstrapping Complete The Look at Pinterest. In *KDD '20: The 26th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, Virtual Event, CA, USA, August 23-27, 2020*, Rajesh Gupta, Yan Liu, Jiliang Tang, and B. Aditya Prakash (Eds.). ACM, 3299–3307. https://doi.org/10.1145/3394486.3403382

[87] Xingchen Li, Xiang Wang, Xiangnan He, Long Chen, Jun Xiao, and Tat-Seng Chua. 2020. Hierarchical Fashion Graph Network for Personalized Outfit Recommendation. In *Proceedings of the 43rd International ACM SIGIR conference on research and development in Information Retrieval, SIGIR 2020, Virtual Event, China, July 25-30, 2020*. ACM, 159–168. https://doi.org/10.1145/3397271.3401080

[88] Yuncheng Li, Liangliang Cao, Jiang Zhu, and Jiebo Luo. 2017. Mining Fashion Outfit Composition Using an End-to-End Deep Learning Approach on Set Data. *IEEE Trans. Multim.* 19, 8 (2017), 1946–1955. https://doi.org/10.1109/TMM.2017.2690144

[89] Lizi Liao, Xiangnan He, Bo Zhao, Chong-Wah Ngo, and Tat-Seng Chua. 2018. Interpretable multimodal retrieval for fashion products. In *Proceedings of the 26th ACM international conference on Multimedia*. 1571–1579.

[90] Yusan Lin, Maryam Moosaei, and Hao Yang. 2020. OutfitNet: Fashion Outfit Recommendation with Attention-Based Multiple Instance Learning. In *WWW '20: The Web Conference 2020, Taipei, Taiwan, April 20-24, 2020*. ACM / IW3C2, 77–87. https://doi.org/10.1145/3366423.3380096

[91] Yujie Lin, Pengjie Ren, Zhumin Chen, Zhaochun Ren, Jun Ma, and Maarten de Rijke. 2019. Improving Outfit Recommendation with Co-supervision of Fashion Generation. In *The World Wide Web Conference, WWW 2019, San Francisco, CA, USA, May 13-17, 2019*. 1095–1105. https://doi.org/10.1145/3308558.3313614

[92] Yujie Lin, Pengjie Ren, Zhumin Chen, Zhaochun Ren, Jun Ma, and Maarten de Rijke. 2020. Explainable Outfit Recommendation with Joint Outfit Matching and Comment Generation. *IEEE Trans. Knowl. Data Eng.* 32, 8 (2020), 1502–1516. https://doi.org/10.1109/TKDE.2019.2906190

[93] Yen-Liang Lin, Son Tran, and Larry S Davis. 2020. Fashion outfit complementary item retrieval. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 3311–3319.

[94] Luyao Liu, Xingzhong Du, Lei Zhu, Fumin Shen, and Zi Huang. 2018. Learning Discrete Hashing Towards Efficient Fashion Recommendation. *Data Science and Engineering* 3, 4 (2018), 307–322. https://doi.org/10.1007/s41019-018-0079-z

[95] Luoqi Liu, Hui Xu, Si Liu, Junliang Xing, Xi Zhou, and Shuicheng Yan. 2013. "Wow! you are so beautiful today!". In *MM '13*.

[96] Qiang Liu, Shu Wu, and Liang Wang. 2017. DeepStyle: Learning User Preferences for Visual Recommendation. In *SIGIR*. ACM, 841–844.

[97] Si Liu, Zheng Song, Guangcan Liu, Changsheng Xu, Hanqing Lu, and Shuicheng Yan. 2012. Street-to-shop: Cross-scenario clothing retrieval via parts alignment and auxiliary set. In *2012 IEEE Conference on Computer Vision and Pattern Recognition*. IEEE, 3330–3337.

[98] Yujie Liu, Yongbiao Gao, Shihe Feng, and Zongmin Li. 2017. Weather-to-garment: Weather-oriented clothing recommendation. In *2017 IEEE International Conference on Multimedia and Expo (ICME)*. IEEE, 181–186.

[99] Ziwei Liu, Ping Luo, Shi Qiu, Xiaogang Wang, and Xiaoou Tang. 2016. Deepfashion: Powering robust clothes recognition and retrieval with rich annotations. In *Proceedings of the IEEE conference on computer vision and pattern recognition*. 1096–1104.

[100] Ling Lo, Chia-Lin Liu, Rong-An Lin, Bo Wu, Hong-Han Shuai, and Wen-Huang Cheng. 2019. Dressing for attention: Outfit based fashion popularity prediction. In *2019 IEEE International Conference on Image Processing (ICIP)*. IEEE, 3222–3226.

[101] Zhi Lu, Yang Hu, Yan Chen, and Bing Zeng. 2021. Personalized Outfit Recommendation With Learnable Anchors. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 12722–12731.

[102] Zhi Lu, Yang Hu, Yunchao Jiang, Yan Chen, and Bing Zeng. 2019. Learning Binary Code for Personalized Fashion Recommendation. In *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2019, Long Beach, CA, USA, June 16-20, 2019*. Computer Vision Foundation / IEEE, 10562–10570. https://doi.org/10.1109/CVPR.2019.01081

[103] Ian MacKenzie, Chris Meyer, and Steve Noble. 2013. How retailers can keep up with consumers. *McKinsey & Company* 18 (2013).

[104] Utkarsh Mall, Kevin Matzen, Bharath Hariharan, Noah Snavely, and Kavita Bala. 2019. GeoStyle: Discovering fashion trends and events. In *ICCV*.

[105] Utkarsh Mall, Kevin Matzen, Bharath Hariharan, Noah Snavely, and Kavita Bala. 2019. GeoStyle: Discovering Fashion Trends and Events. In *2019 IEEE/CVF International Conference on Computer Vision, ICCV 2019, Seoul, Korea (South), October 27 - November 2, 2019*. IEEE, 411–420. https://doi.org/10.1109/ICCV.2019.00050

[106] Kevin Matzen, Kavita Bala, and Noah Snavely. 2017. StreetStyle: Exploring world-wide clothing styles from millions of photos. *arXiv preprint arXiv:1706.01869* (2017).

[107] Julian J. McAuley, Christopher Targett, Qinfeng Shi, and Anton van den Hengel. 2015. Image-Based Recommendations on Styles and Substitutes. In *Proceedings of the 38th International ACM SIGIR Conference on Research and Development in Information Retrieval, Santiago, Chile, August 9-13, 2015*. 43–52. https://doi.org/10.1145/2766462.2767755

[108] Augustus Odena, Christopher Olah, and Jonathon Shlens. 2017. Conditional Image Synthesis with Auxiliary Classifier GANs. In *ICML (Proceedings of Machine Learning Research)*, Vol. 70. PMLR, 2642–2651.

[109] Ioannis Pachoulakis and Kostas Kapetanakis. 2012. Augmented reality platforms for virtual fitting rooms. *The International Journal of Multimedia & Its Applications* 4, 4 (2012), 35.

[110] Charles Packer, Julian McAuley, and Arnau Ramisa. 2018. Visually-aware personalized recommendation using interpretable image representations. *Proceedings of the third international workshop on fashion and KDD (held in conjunction with KDD'18)* (2018).

[111] Tse-Yu Pan, Yi-Zhu Dai, Min-Chun Hu, and Wen-Huang Cheng. 2019. Furniture style compatibility recommendation with cross-class triplet loss. *Multim. Tools Appl.* 78, 3 (2019), 2645–2665. https://doi.org/10.1007/s11042-018-5747-5

[112] Jacqueline Parr and Sanjukta Pookulangara. 2017. The impact of True Fit Technology on Consumer Confidence in their online clothing purchase. In *International Textile and Apparel Association Annual Conference Proceedings*, Vol. 74. Iowa State University Digital Press.

[113] Luisa F. Polanía and Satyajit Gupte. 2019. Learning Fashion Compatibility Across Apparel Categories for Outfit Recommendation. In *2019 IEEE International Conference on Image Processing, ICIP 2019, Taipei, Taiwan, September 22-25, 2019*. 4489–4493. https://doi.org/10.1109/ICIP.2019.8803587

[114] Xueming Qian, He Feng, Guoshuai Zhao, and Tao Mei. 2014. Personalized Recommendation Combining User Interest and Social Circle. *IEEE Trans. Knowl. Data Eng.* 26, 7 (2014), 1763–1777. https://doi.org/10.1109/TKDE.2013.168

[115] Yu Qian, P. Giaccone, Michele Sasdelli, Eduard Vasquez, and Biswa Sengupta. 2017. Algorithmic clothing: hybrid recommendation, from street-style-to-shop. *CoRR* abs/1705.09451 (2017). arXiv:1705.09451 http://arxiv.org/abs/1705.09451

[116] Ruihong Qiu, Sen Wang, Zhi Chen, Hongzhi Yin, and Zi Huang. 2021. CausalRec: Causal Inference for Visual Debiasing in Visually-Aware Recommendation. *CoRR* abs/2107.02390 (2021). arXiv:2107.02390 https://arxiv.org/abs/2107.02390

[117] Pei-Luen Patrick Rau, Ziyang Li, and Dinglong Huang. 2020. Who Should Provide Clothing Recommendation Services: Artificial Intelligence or Human Experts? *J. Inf. Technol. Res.* 13, 3 (jul 2020), 113–125. https://doi.org/10.4018/JITR.2020070107

[118] Abhinav Ravi, Sandeep Repakula, Ujjal Kr Dutta, and Maulik Parmar. 2020. Buy Me That Look: An Approach for Recommending Similar Fashion Products. *CoRR* abs/2008.11638 (2020). arXiv:2008.11638 https://arxiv.org/abs/2008.11638

[119] Steffen Rendle, Christoph Freudenthaler, Zeno Gantner, and Lars Schmidt-Thieme. 2009. BPR: Bayesian Personalized Ranking from Implicit Feedback. In *UAI*. AUAI Press, 452–461.

[120] Steffen Rendle, Christoph Freudenthaler, Zeno Gantner, and Lars Schmidt-Thieme. 2009. BPR: Bayesian personalized ranking from implicit feedback. In *Proceedings of the twenty-fifth conference on uncertainty in artificial intelligence*. AUAI Press, 452–461.

[121] Tim Salimans, Ian J. Goodfellow, Wojciech Zaremba, Vicki Cheung, Alec Radford, and Xi Chen. 2016. Improved Techniques for Training GANs. In *NIPS*. 2226–2234.

[122] Hosnieh Sattar, Gerard Pons-Moll, and Mario Fritz. 2019. Fashion is taking shape: Understanding clothing preference based on body shape from online sources. In *2019 IEEE Winter Conference on Applications of Computer Vision (WACV)*. IEEE, 968–977.

[123] Dandan Sha, Daling Wang, Xiangmin Zhou, Shi Feng, Yifei Zhang, and Ge Yu. 2016. An Approach for Clothing Recommendation Based on Multiple Image Attributes. In *Web-Age Information Management - 17th International Conference, WAIM 2016, Nanchang, China, June 3-5, 2016, Proceedings, Part I (Lecture Notes in Computer Science)*, Vol. 9658. Springer, 272–285. https://doi.org/10.1007/978-3-319-39937-9_21

[124] Abdul-Saboor Sheikh, Romain Guigourès, Evgenii Koriagin, Yuen King Ho, Reza Shirvany, Roland Vollgraf, and Urs Bergmann. 2019. A deep learning system for predicting size and fit in fashion e-commerce. In *Proceedings of the 13th ACM conference on recommender systems*. 110–118.

[125] Yue Shi, Martha Larson, and Alan Hanjalic. 2014. Collaborative filtering beyond the user-item matrix: A survey of the state of the art and future challenges. *ACM Computing Surveys (CSUR)* 47, 1 (2014), 3.

[126] Yong-Siang Shih, Kai-Yueh Chang, Hsuan-Tien Lin, and Min Sun. 2018. Compatibility Family Learning for Item Recommendation and Generation. In *AAAI*. AAAI Press, 2403–2410.

[127] Edgar Simo-Serra, Sanja Fidler, Francesc Moreno-Noguer, and Raquel Urtasun. 2015. Neuroaesthetics in fashion: Modeling the perception of fashionability. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 869–877.

[128] Montek Singh, Utkarsh Bajpai, Vijayarajan V., and Surya Prasath. 2020. Generation of fashionable clothes using generative adversarial networks: A preliminary feasibility study. *International Journal of Clothing Science and Technology* 32 (04 2020), 177–187. https://doi.org/10.1108/IJCST-12-2018-0148

[129] Xuemeng Song, Fuli Feng, Jinhuan Liu, Zekun Li, Liqiang Nie, and Jun Ma. 2017. NeuroStylist: Neural Compatibility Modeling for Clothing Matching. In *Proceedings of the 25th ACM International Conference on Multimedia* (Mountain View, California, USA) *(MM '17)*. Association for Computing Machinery, New York, NY, USA, 753–761. https://doi.org/10.1145/3123266.3123314

[130] M Sridevi, N ManikyaArun, M Sheshikala, and E Sudarshan. 2020. Personalized fashion recommender system with image based neural networks. *IOP Conference Series: Materials Science and Engineering* 981 (12 2020), 022073. https://doi.org/10.1088/1757-899X/981/2/022073

[131] Guang-Lu Sun, Zhi-Qi Cheng, Xiao Wu, and Qiang Peng. 2018. Personalized clothing recommendation combining user social circle and fashion style consistency. *Multimedia Tools Appl.* 77, 14 (2018), 17731–17754. https://doi.org/10.1007/s11042-017-5245-1

[132] Christian Szegedy, Sergey Ioffe, Vincent Vanhoucke, and Alexander A. Alemi. 2017. Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning. In *Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence, February 4-9, 2017, San Francisco, California, USA*. 4278–4284. http://aaai.org/ocs/index.php/AAAI/AAAI17/paper/view/14806

[133] Pongsate Tangseng and Takayuki Okatani. 2020. Toward Explainable Fashion Recommendation. In *IEEE Winter Conference on Applications of Computer Vision, WACV 2020, Snowmass Village, CO, USA, March 1-5, 2020*. IEEE, 2142–2151. https://doi.org/10.1109/WACV45572.2020.9093367

[134] Pongsate Tangseng and Takayuki Okatani. 2020. Toward explainable fashion recommendation. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*. 2153–2162.

[135] Marika Tiggemann, Amy Slater, Belinda Bury, Kimberley Hawkins, and Bonny Firth. 2013. Disclaimer labels on fashion magazine advertisements: Effects on social comparison and body dissatisfaction. *Body Image* 10, 1 (2013), 45–53. https://doi.org/10.1016/j.bodyim.2012.08.001

[136] Nava Tintarev and Judith Masthoff. 2015. Explaining recommendations: Design and evaluation. In *Recommender systems handbook*. Springer, 353–382.

[137] Hessel Tuinhof, Clemens Pirker, and Markus Haltmeier. 2018. Image-Based Fashion Product Recommendation with Deep Learning. In *Machine Learning, Optimization, and Data Science - 4th International Conference, LOD 2018, Volterra, Italy, September 13-16, 2018, Revised Selected Papers (Lecture Notes in Computer Science)*, Vol. 11331. Springer, 472–481. https://doi.org/10.1007/978-3-030-13709-0_40

[138] Laura Vandenbosch and Steven Eggermont. 2012. Understanding Sexual Objectification: A Comprehensive Approach Toward Media Exposure and Girls' Internalization of Beauty Ideals, Self-Objectification, and Body Surveillance. *Journal of Communication* 62, 5 (08 2012), 869–887. https://doi.org/10.1111/j.1460-2466.2012.01667.x arXiv:https://academic.oup.com/joc/article-pdf/62/5/869/22321787/jjnlcom0869.pdf

[139] Andreas Veit, Balazs Kovacs, Sean Bell, Julian J. McAuley, Kavita Bala, and Serge J. Belongie. 2015. Learning Visual Clothing Style with Heterogeneous Dyadic Co-Occurrences. In *2015 IEEE International Conference on Computer Vision, ICCV 2015, Santiago, Chile, December 7-13, 2015*. IEEE Computer Society, 4642–4650. https://doi.org/10.1109/ICCV.2015.527

[140] Sirion Vittayakorn, Kota Yamaguchi, Alexander C Berg, and Tamara L Berg. 2015. Runway to realway: Visual analysis of fashion. In *2015 IEEE Winter Conference on Applications of Computer Vision*. IEEE, 951–958.

[141] Riccardo Volpi, Pietro Morerio, Silvio Savarese, and Vittorio Murino. 2018. Adversarial Feature Augmentation for Unsupervised Domain Adaptation. In *2018 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2018, Salt Lake City, UT, USA, June 18-22, 2018*. IEEE Computer Society, 5495–5504. https://doi.org/10.1109/CVPR.2018.00576

[142] Y. Wakita, Kenta Oku, and K. Kawagoe. 2016. Toward Fashion-Brand Recommendation Systems Using Deep-Learning: Preliminary Analysis.

[143] Kaili Wang, Yu-Hui Huang, Jose Oramas, Luc Van Gool, and Tinne Tuytelaars. 2018. An analysis of human-centered geolocation. In *2018 IEEE Winter Conference on Applications of Computer Vision (WACV)*. IEEE, 2058–2066.

[144] Lichuan Wang, Xianyi Zeng, Ludovic Koehl, and Yan Chen. 2015. Intelligent Fashion Recommender System: Fuzzy Logic in Personalized Garment Design. *IEEE Trans. Human-Machine Systems* 45, 1 (2015), 95–109.

[145] Wenguan Wang, Yuanlu Xu, Jianbing Shen, and Song-Chun Zhu. 2018. Attentive fashion grammar network for fashion landmark detection and clothing category classification. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 4271–4280.

[146] Haokun Wen, Xuemeng Song, Xin Yang, Yibing Zhan, and Liqiang Nie. 2021. Comprehensive Linguistic-Visual Composition Network for Image Retrieval. In *SIGIR '21: The 44th International ACM SIGIR Conference on Research and Development in Information Retrieval, Virtual Event, Canada, July 11-15, 2021*, Fernando Diaz, Chirag Shah, Torsten Suel, Pablo Castells, Rosie Jones, and Tetsuya Sakai (Eds.). ACM, 1369–1378. https://doi.org/10.1145/3404835.3462967

[147] Yufan Wen, Xiaoqiang Liu, and Bo Xu. 2018. Personalized Clothing Recommendation Based on Knowledge Graph. In *2018 International Conference on Audio, Language and Image Processing (ICALIP)*. 1–5. https://doi.org/10.1109/ICALIP.2018.8455311

[148] Matthias Wölbitsch, Thomas Hasler, Simon Walk, and Denis Helic. 2020. Mind the Gap: Exploring Shopping Preferences Across Fashion Retail Channels. In *Proceedings of the 28th ACM Conference on User Modeling, Adaptation and Personalization, UMAP 2020, Genoa, Italy, July 12-18, 2020*, Tsvi Kuflik, Ilaria Torre, Robin Burke, and Cristina Gena (Eds.). ACM, 257–265. https://doi.org/10.1145/3340631.3394866

[149] Qianqian Wu, Pengpeng Zhao, and Zhiming Cui. 2020. Visual and textual jointly enhanced interpretable fashion recommendation. *IEEE Access* 8 (2020), 68736–68746.

[150] Hao Xing, Zhike Han, and Yichen Shen. 2020. ClothNet: A Neural Network Based Recommender System. In *Fuzzy Systems and Data Mining VI - Proceedings of FSDM 2020, Virtual Event, November 13-16, 2020 (Frontiers in Artificial Intelligence and Applications)*, Vol. 331. IOS Press, 261–272. https://doi.org/10.3233/FAIA200706

[151] Kota Yamaguchi, M Hadi Kiapour, Luis E Ortiz, and Tamara L Berg. 2012. Parsing clothing in fashion photographs. In *2012 IEEE Conference on Computer vision and pattern recognition*. IEEE, 3570–3577.

[152] Cairong Yan, Yizhou Chen, and Lingjie Zhou. 2019. Differentiated Fashion Recommendation Using Knowledge Graph and Data Augmentation. *IEEE Access* 7 (2019), 102239–102248. https://doi.org/10.1109/ACCESS.2019.2928848

[153] Ming Yang and Kai Yu. 2011. Real-time clothing recognition in surveillance videos. In *2011 18th IEEE International Conference on Image Processing*. IEEE, 2937–2940.

[154] Wei Yang, Masahiro Toyoura, and Xiaoyang Mao. 2012. Hairstyle Suggestion Using Statistical Learning, Vol. 7131. 277–287. https://doi.org/10.1007/978-3-642-27355-1_27

[155] Xun Yang, Xiangnan He, Xiang Wang, Yunshan Ma, Fuli Feng, Meng Wang, and Tat-Seng Chua. 2019. Interpretable fashion matching with rich attributes. In *Proceedings of the 42nd International ACM SIGIR Conference on Research and Development in Information Retrieval*. 775–784.

[156] Zilin Yang, Zhuo Su, Yang Yang, and Ge Lin. 2018. From recommendation to generation: A novel fashion clothing advising framework. In *2018 7th International Conference on Digital Home (ICDH)*. IEEE, 180–186.

[157] Zilin Yang, Zhuo Su, Yang Yang, and Ge Lin. 2018. From Recommendation to Generation: A Novel Fashion Clothing Advising Framework. *2018 7th International Conference on Digital Home (ICDH)* 1, 1 (2018), 180–186.

[158] Ruiping Yin, Kan Li, Jie Lu, and Guangquan Zhang. 2019. Enhancing Fashion Recommendation with Visual Compatibility Relationship. In *The World Wide Web Conference, WWW 2019, San Francisco, CA, USA, May 13-17, 2019*. ACM, 3434–3440. https://doi.org/10.1145/3308558.3313739

[159] Tong Yu, Yilin Shen, and Hongxia Jin. 2019. A visual dialog augmented interactive recommender system. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*. 157–165.

[160] Wenhui Yu, Xiangnan He, Jian Pei, Xu Chen, Li Xiong, Jinfei Liu, and Zheng Qin. 2021. Visually aware recommendation with aesthetic features. *VLDB J.* 30, 4 (2021), 495–513. https://doi.org/10.1007/s00778-021-00651-y

[161] Huijing Zhan, Jie Lin, Kenan Emir Ak, Boxin Shi, Ling-Yu Duan, and Alex C. Kot. 2021. A3-FKG: Attentive Attribute-Aware Fashion Knowledge Graph for Outfit Preference Prediction. *IEEE Transactions on Multimedia* (2021), 1–1. https://doi.org/10.1109/TMM.2021.3059514

[162] Xishan Zhang, Jia Jia, Ke Gao, Yongdong Zhang, Dongming Zhang, Jintao Li, and Qi Tian. 2017. Trip Outfits Advisor: Location-Oriented Clothing Recommendation. *IEEE Transactions on Multimedia* 19, 11 (2017), 2533–2544. https://doi.org/10.1109/TMM.2017.2696825

[163] Yan Zhang, Xiang Liu, Yunyu Shi, Yunqi Guo, Chaoqun Xu, Erwen Zhang, Jiaxun Tang, and Zhijun Fang. 2017. Fashion Evaluation Method for Clothing Recommendation Based on Weak Appearance Feature. *Scientific Programming* 2017 (2017), 8093057:1–8093057:12.

[164] Kui Zhao, Xia Hu, Jiajun Bu, and Can Wang. 2017. Deep Style Match for Complementary Recommendation. In *The Workshops of the The Thirty-First AAAI Conference on Artificial Intelligence, Saturday, February 4-9, 2017, San Francisco, California, USA (AAAI Workshops)*, Vol. WS-17. AAAI Press. http://aaai.org/ocs/index.php/WS/AAAIW17/paper/

view/15069

[165] Jiachen Zhong, Rita Tse, Gustavo Marfia, and Giovanni Pau. 2019. Fashion popularity analysis based on online social network via deep learning. In *Eleventh International Conference on Digital Image Processing (ICDIP 2019)*, Vol. 11179. International Society for Optics and Photonics, 1117938.

[166] Wei Zhou, P. Y. Mok, Yanghong Zhou, Yangping Zhou, Jialie Shen, Qiang Qu, and K. P. Chau. 2019. Fashion recommendations through cross-media information retrieval. *J. Visual Communication and Image Representation* 61 (2019), 112–120.

[167] Zhengzhong Zhou, Xiu Di, Wei Zhou, and Liqing Zhang. 2018. Fashion Sensitive Clothing Recommendation Using Hierarchical Collocation Model. In *ACM Multimedia*. ACM, 1119–1127.

[168] Shizhan Zhu, Sanja Fidler, Raquel Urtasun, Dahua Lin, and Chen Change Loy. 2017. Be Your Own Prada: Fashion Synthesis with Structural Coherence. In *IEEE International Conference on Computer Vision, ICCV 2017, Venice, Italy, October 22-29, 2017*. IEEE Computer Society, 1689–1697. https://doi.org/10.1109/ICCV.2017.186

[169] Kevin Zielnicki. 2019. Simulacra and Selection: Clothing Set Recommendation at Stitch Fix. In *Proceedings of the 42nd International ACM SIGIR Conference on Research and Development in Information Retrieval*. 1379–1380.

[170] Susana Zoghbi, Geert Heyman, Juan Carlos Gomez, and Marie-Francine Moens. 2016. Fashion meets computer vision and nlp at e-commerce search. *International Journal of Computer and Electrical Engineering* 8, 1 (2016), 31–43.

[171] Qin Zou, Zheng Zhang, Qian Wang, Qingquan Li, Long Chen, and Song Wang. 2016. Who Leads the Clothing Fashion: Style, Color, or Texture? A Computational Study. *CoRR* abs/1608.07444 (2016). arXiv:1608.07444 http://arxiv.org/abs/1608.07444